



*influx*/days

# Basic Data Analysis

Emanuele Della Valle

Prof. @ Politecnico di Milano

Founder & Partner @ Quantia Consulting

Marco Balduini

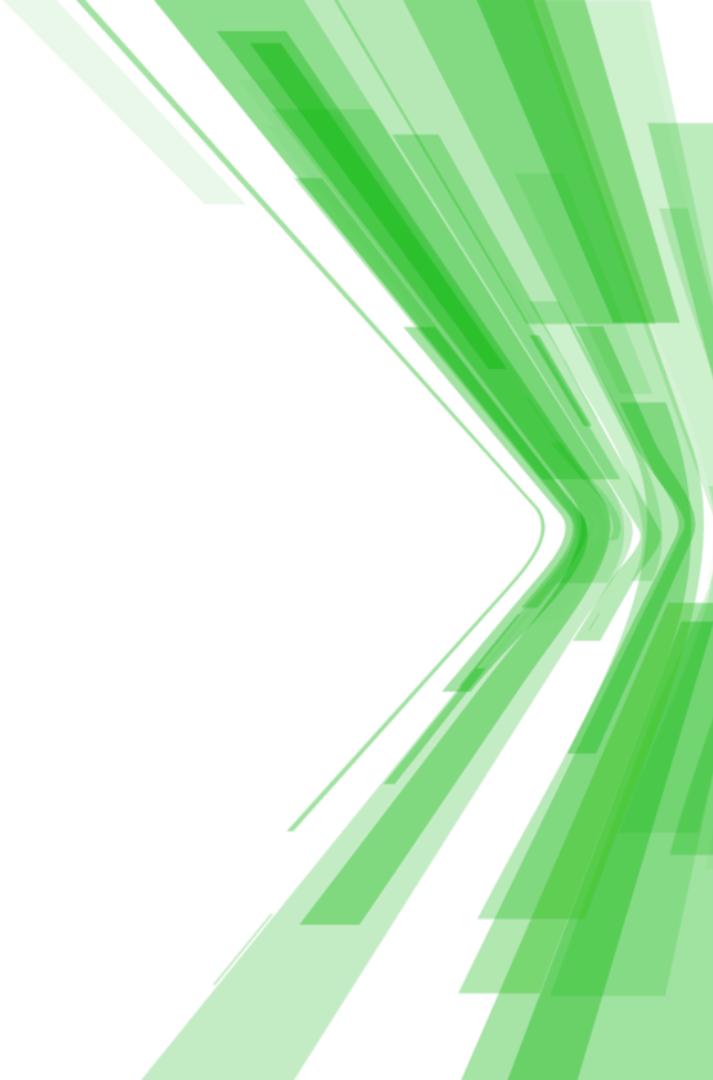
Founder & CEO @ Quantia Consulting

Riccardo Tommasini

Prof. @ University of Tartu



# Introduction

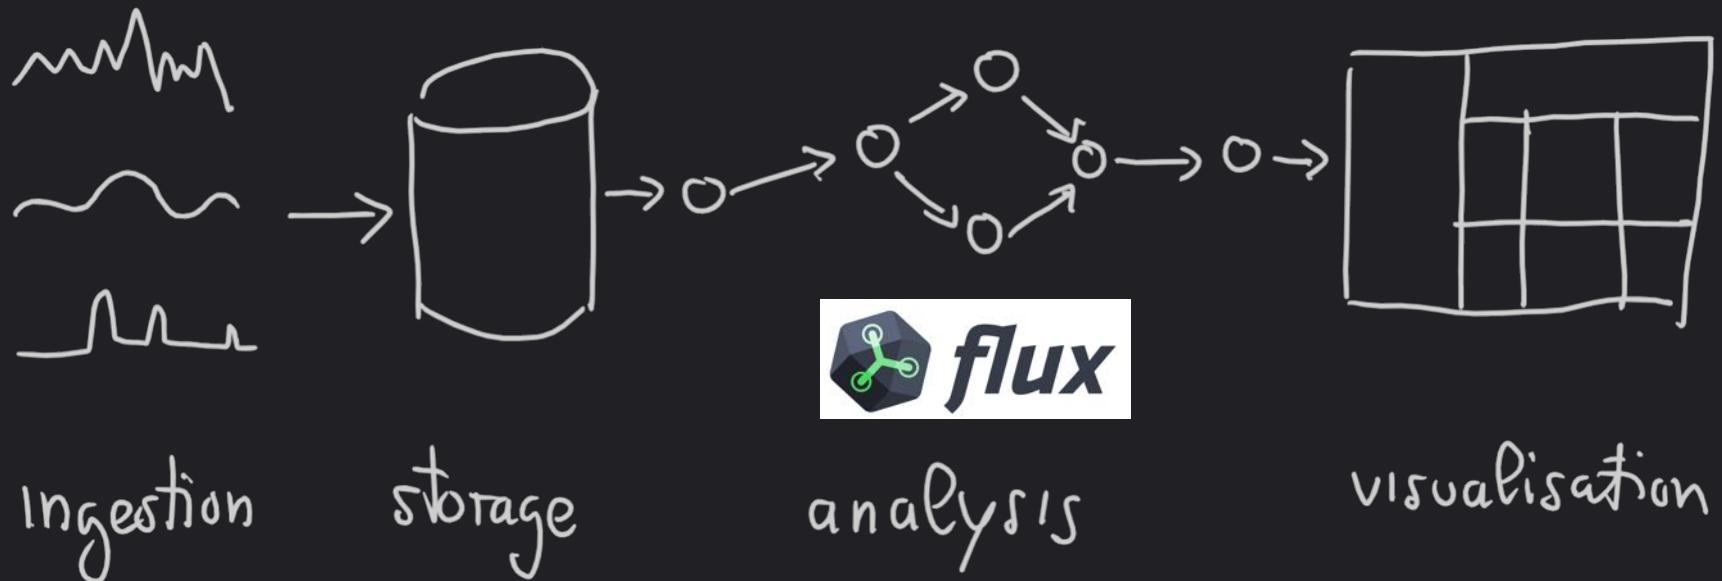


# Data Lifecycle



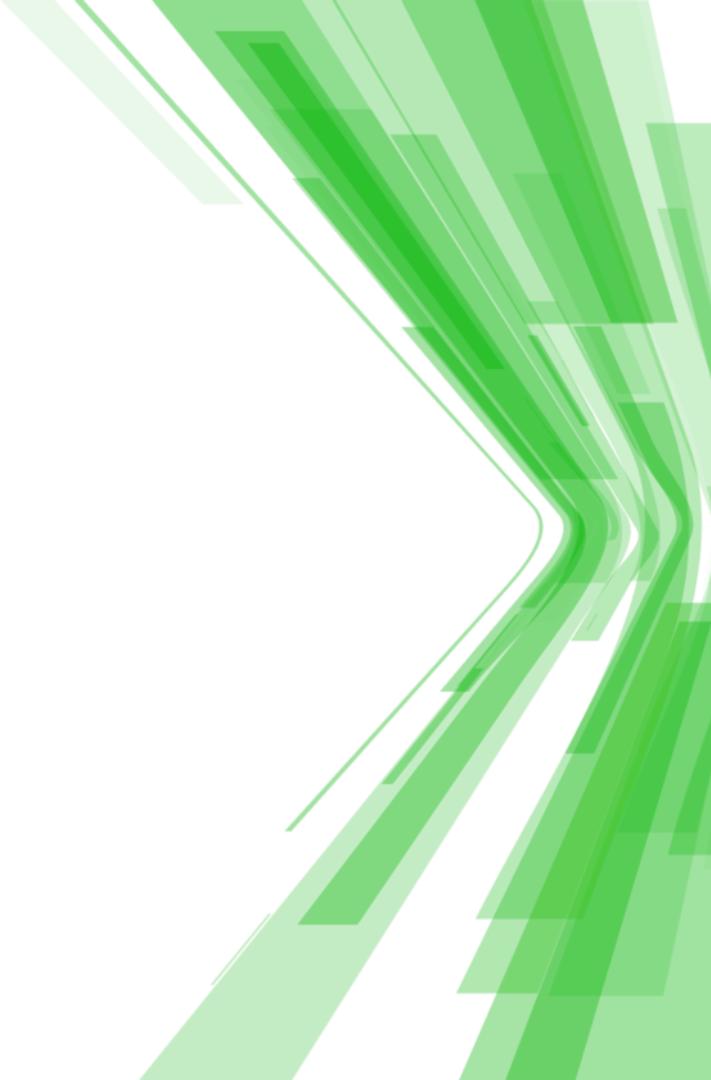
ingestion      storage

# Data Lifecycle





*flux*



# Why Flux?

I don't want to have to write query code in one language and processing code in another language.

– Paul Dix - CTO InfluxData

# What Is Flux?

Flux ~~Query~~ Language?

Query + Process + Act

Flux Data Scripting Language

# What Is Flux?

- Written to be:
  - Useable: easy to learn
  - Readable: developers read more code than we write
  - Extensible: developers can extend the language
  - Composable: developers can build onto the language
  - Testable: queries are code
  - Contributable: open source contributions matter

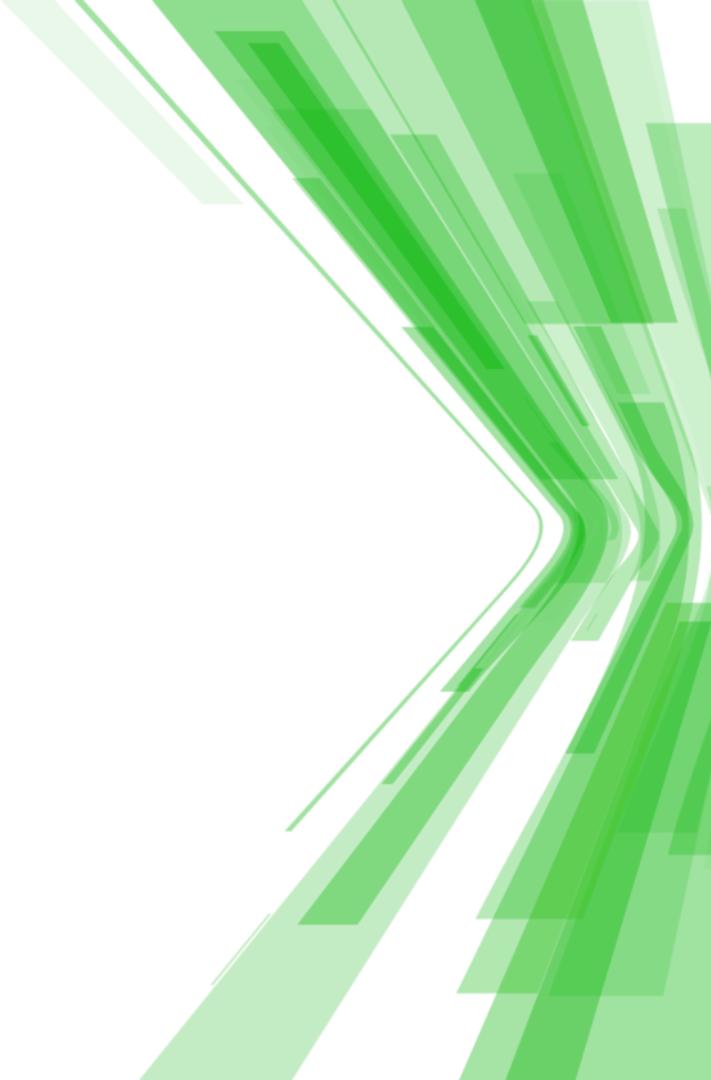
## What is Flux? (cont.)

- Open Source (MIT)
- Data-first mentality
- Streaming Data
- Planner-optimized
- Compatible with multiple sources and sinks

## Flux vs. InfluxQL

- Flux is an alternative to InfluxQL with additional functionality like:
  - Joins
  - Math across measurements
  - Sort by any column
  - Pivot
  - Histograms
  - Covariance
- Reference: <https://v2.docs.influxdata.com/v2.0/query-data/>

# Exploring flux Basics



# Basic Flux Language Elements

- Anatomy of a Flux query:
  - data source
  - time range
  - data filters

**Pipe-forward operator:**  
chain operations together

```
from(bucket:"telegraf/autogen")  
|> range(start:-1h)  
|> filter(fn: (r) =>  
    r._measurement == "cpu" and  
    r._field == "system" and  
    r._value > 2000.0)
```

**Anonymous function**

# Basic Flux Language Elements

- Anatomy of a Flux query:

- data source
- time range
- data filters

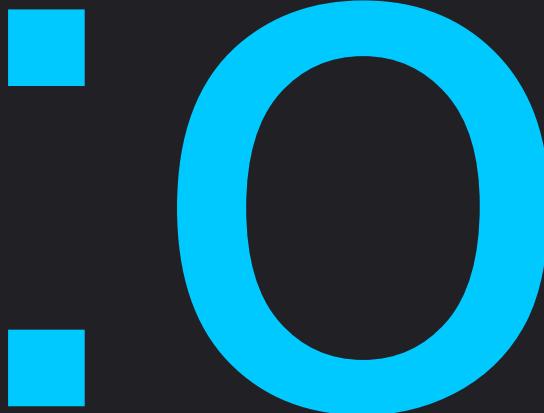
**expression(s)**

```
from(bucket:"telegraf/autogen")  
|> range(start:-1h)  
|> filter(fn: (r) =>  
    {  
        r._measurement == "cpu" and  
        r._field == "system" and  
        r._value > 2000.0  
    })
```

**parameter(s)**

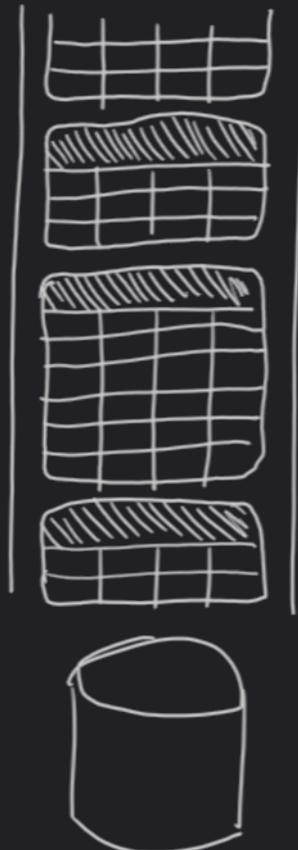
# Flux Data Model

An infinite stream...  
of finite tables...  
identified by the GroupKey



# Flux Data Model

An infinite stream...  
of finite tables...  
identified by the GroupKey



# From bucket to tables

_time	_m	host	_field	_value	_time	_m	host	_field	_value	_time	_m	host	_field	_value
14...0	cpu	serverA	system	2342	14...1	cpu	serverB	system	2779	14...1	cpu	serverB	idle	1789
14...2	cpu	serverA	system	477	14...3	cpu	serverB	system	1274	14...2	cpu	serverB	idle	2475
GroupKey[cpu, serverA, system]					GroupKey[cpu, serverB, system]					GroupKey[cpu, serverB, idle]				



```
from(bucket:"telegraf/autogen")
|> range(start:-1s)
```

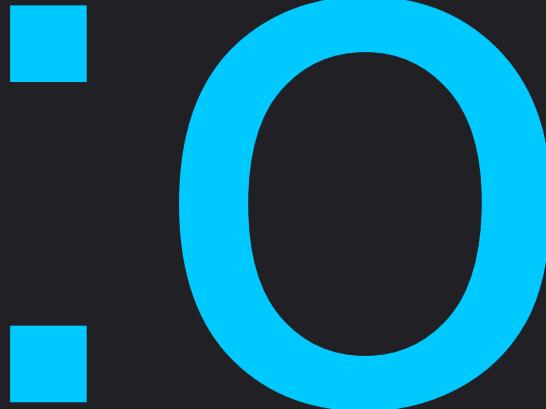
_time	_m	host	idle	system
14...0	cpu	serverA	1667	2342
...	...	...	...	...
14...1	cpu	serverB	1789	2779

\_m = \_measurement

# Query Process

For Each Table:

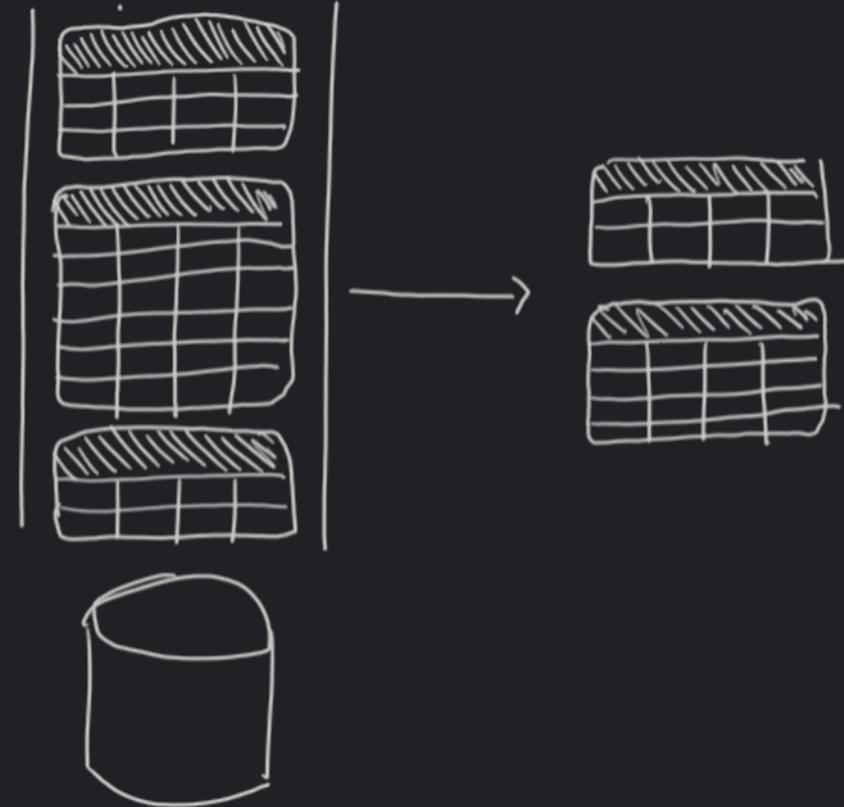
1. Convert incoming rows into  
0 or more outgoing rows
2. Map outgoing rows to  
tables



# Query Process

For Each Table:

1. Convert incoming rows into 0 or more outgoing rows
2. Map outgoing rows to tables



_time	_m	host	_field	_value
14...1	cpu	serverB	idle	1789
14...2	cpu	serverB	idle	2475

GroupKey[cpu, serverB, idle]

_time	_m	host	_field	_value
14...1	cpu	serverB	system	2779
14...3	cpu	serverB	system	1274

GroupKey[cpu, serverB, system]

_time	_m	host	_field	_value
14...0	cpu	serverA	system	2342
14...2	cpu	serverA	system	477

GroupKey[cpu, serverA, system]

```
|> filter(fn: (r) =>
  r._measurement == "cpu" and
  r._field == "system" and
  r._value > 2000.0)
```



∅

_time	_m	host	_field	_value
14...1	cpu	serverB	system	2779

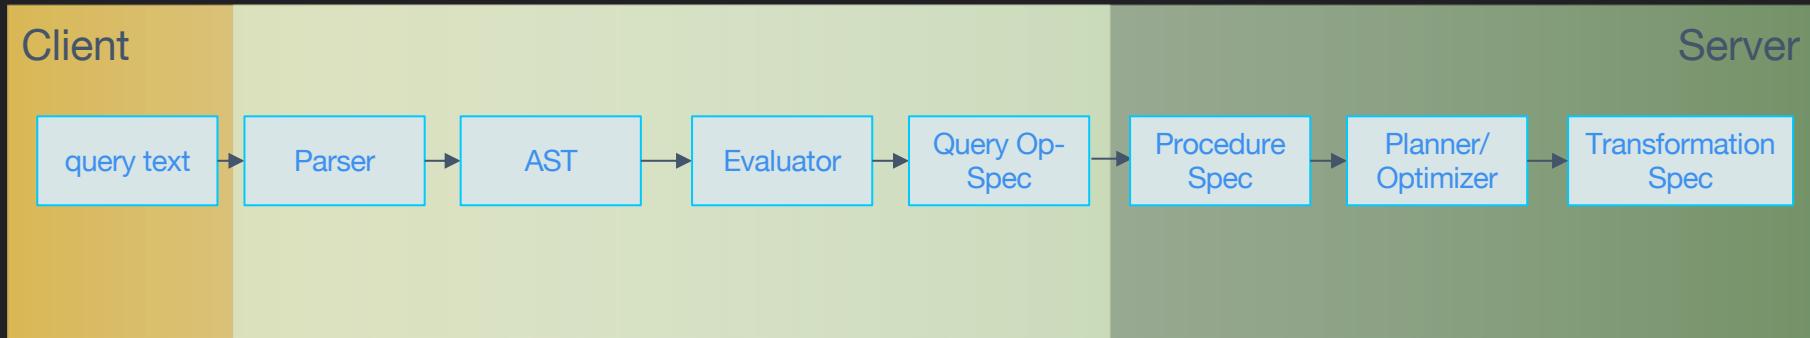
GroupKey[cpu, serverB, system]

[...]

_time	_m	host	_field	_value
14...0	cpu	serverA	system	2342

GroupKey[cpu, serverA, system]

# The Life of a Flux Query



# Quiz

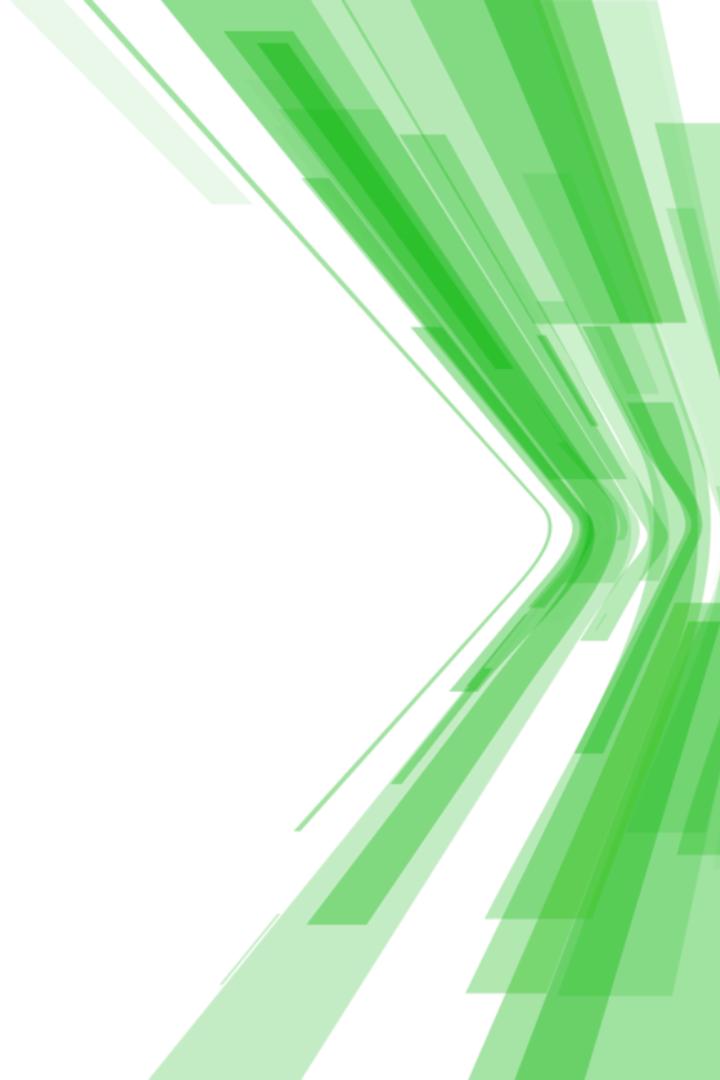
1. What's a Bucket?  
A. List of columns for which every row in a table has the same value
2. What's the Pipe-forward operator `|>` for?  
B. A named data source with retention policy
3. What's a Table in Flux?  
C. Chains operations
4. What's a Group Key in Flux?  
D. A part of a result

## Quiz answer

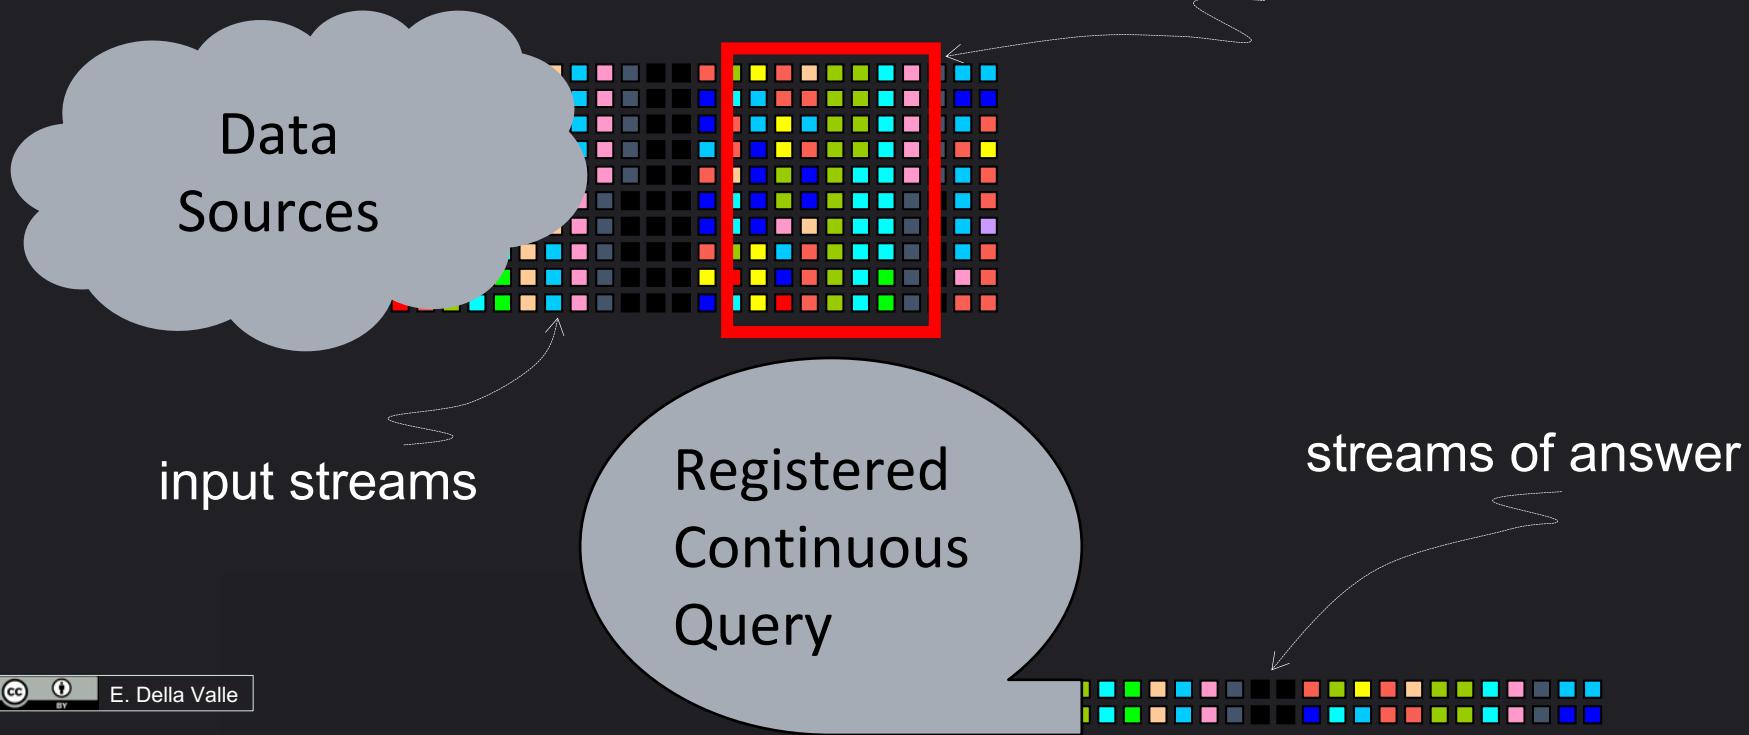
1. What's a Bucket?  
A. List of columns for which every row in a table has the same value
2. What's the Pipe-forward operator `|>` for?  
B. A named data source with retention policy
3. What's a Table in Flux?  
C. Chains operations
4. What's a Group Key in Flux?  
D. A part of a result

[1:B, 2:C, 3:D, 4:A]

# Exploring flux Windowing



# Key concepts



# Windows

- Streams are infinite in nature
- Processing can end only if the input is finite
- A **window operator** selects a finite portion of a stream
  - **Time range** is a window operator whose input is a time interval specified either stating an **absolute** start and end time

```
range(start: 2019-10-01T00:00:00Z,  
      stop: 2019-10-01T00:05:00Z)
```
  - or stating a **relative** time period relative

```
range(start:-2h, stop: -1h)  
range(start:-1h)
```

Let's get dirty!

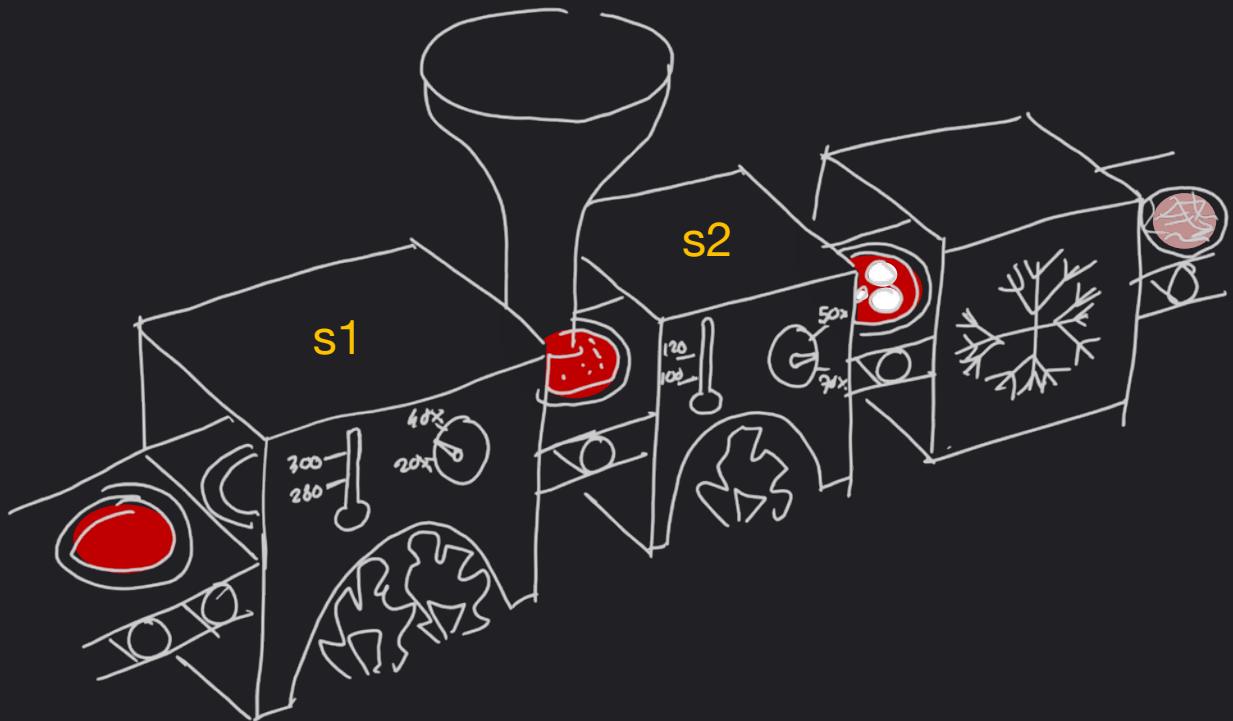


2

# Continuous Linear Pizza Oven

## Learning goals:

- Data source
- Window
- Tables



# Task

Extract all the measurements in a given range

- Absolute: from 2020-04-30T10:00:00Z to 2020-04-30T10:05:00Z
- Relative: in the last 24 hours

## Take home message

The **range()** function **extracts** from the named bucket all the possible **tables in the given time range**

# Let's get dirty!

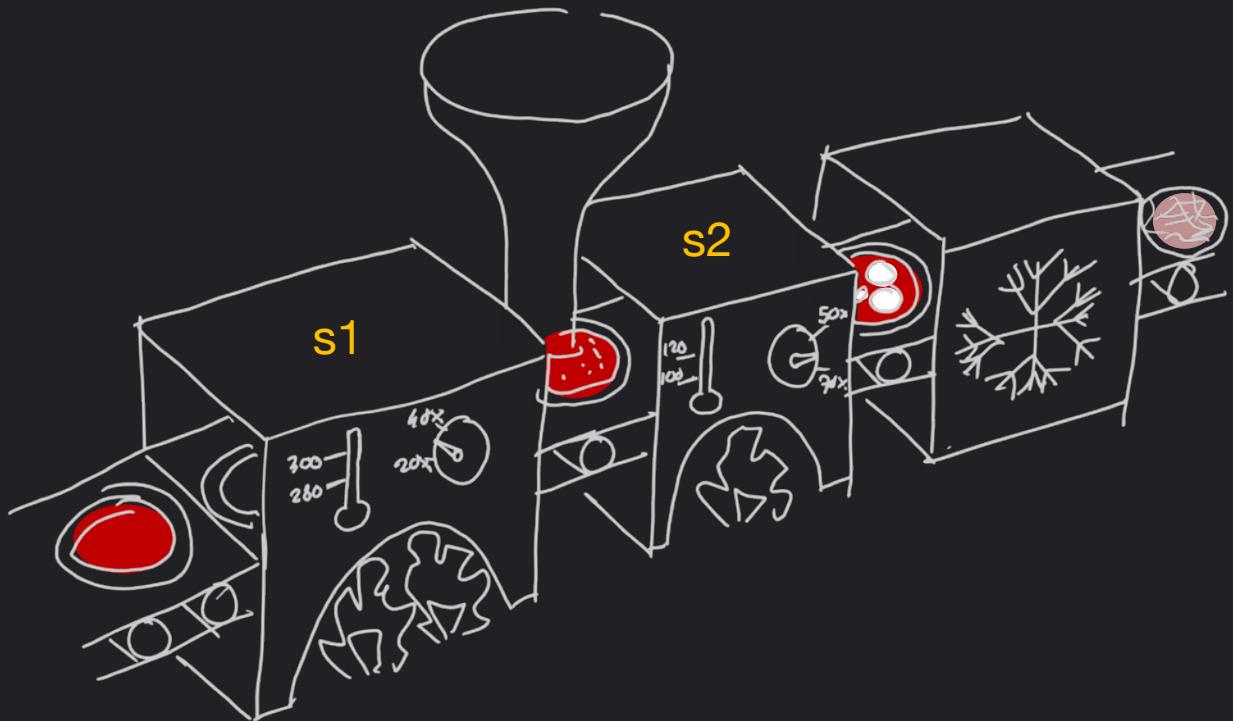


3

# Continuous Linear Pizza Oven

## Learning goals:

- Filter by tag



## Task

Extract the temperature data from the cooking base area  
(sensor S1)

## Take home message

The **filter()** function **applied to tags selects** only the **table(s)** of your interest from all the possible tables

Let's get dirty!

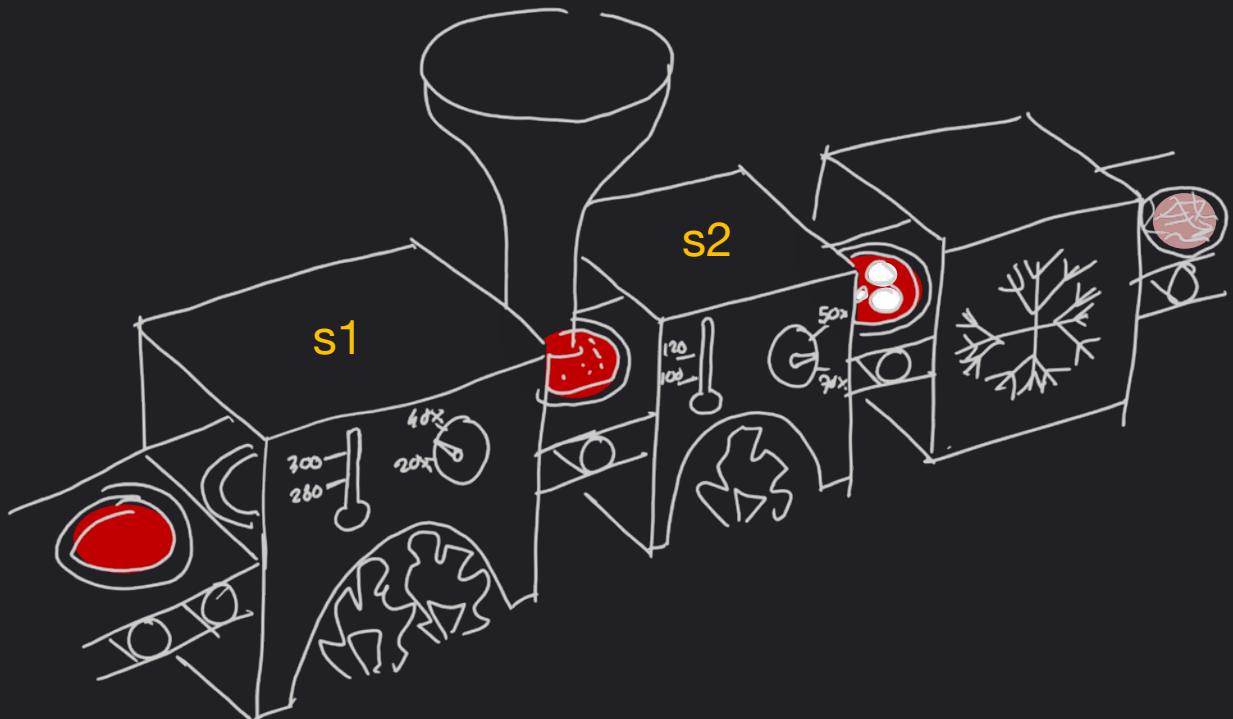


4

# Continuous Linear Pizza Oven

Learning goals:

- Filter by value



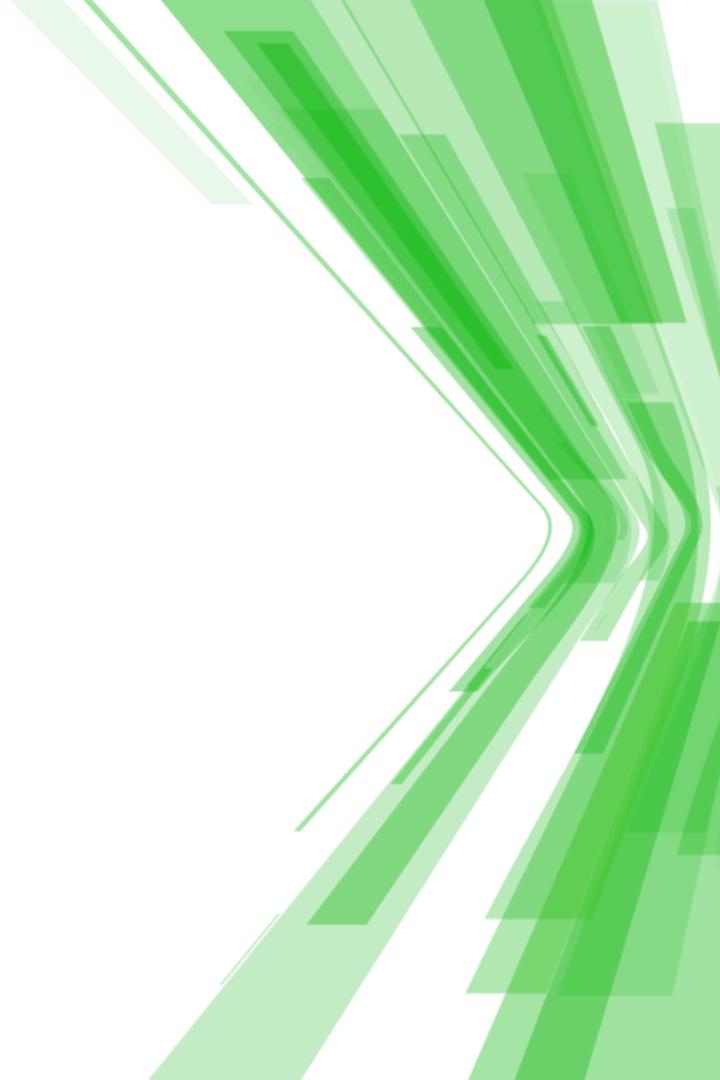
## Task

Extract the measurements from the cooking base area (sensor S1) with a temperature under 300°

## Take home message

The `filter()` function **applied to values selects** only the **row(s)** of your interest from all the possible tables

# Exploring flux Shaping and Processing

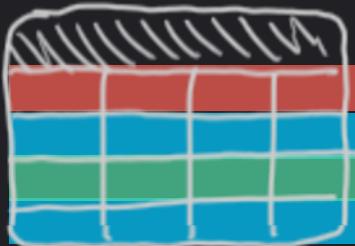
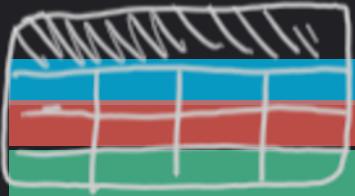


# All Flux Language Elements

- Anatomy of a Flux query:

- data source                         `from(bucket:"telegraf/autogen")`
- time range                         `|> range(start:-1h)`
- data filters                         `|> filter(fn: (r) =>`
- shaping                                 `r._measurement == "cpu" and`
- processing                             `r._field == "system" and`  
   `r._value > 2000.0)`

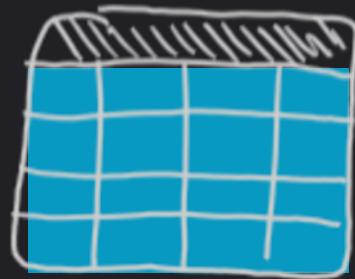
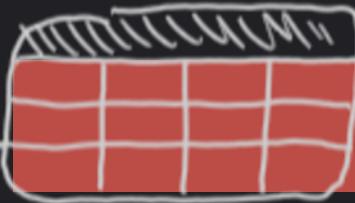
- processing                             `|> group(columns: ["_field"])`
- processing                             `|> mean()`



shape  
→



process  
→





Exploring flux  
Shaping with group  
Processing with aggregators and  
selectors

_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:45	cpu	serverA	system	3000
GroupKey[cpu, serverA, system]				

_time	_m	host	_field	_value
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
GroupKey[cpu, serverB, system]				

_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
GroupKey[cpu, serverA, idle]				

_time	_m	host	_field	_value
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400
GroupKey[cpu, serverB, idle]				

group(columns: ["\_field"])



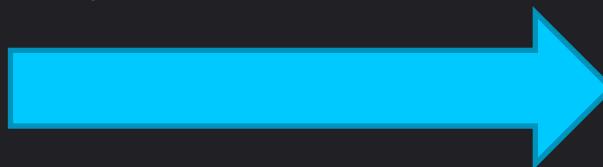
_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
11:45	cpu	serverA	system	3000
GroupKey[cpu, system]				

_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400
GroupKey[cpu, idle]				

# Example of aggregator

_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
11:45	cpu	serverA	system	3000
GroupKey[cpu, system]				

| > mean()



_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400
GroupKey[cpu, idle]				

_m	_field	_value
cpu	system	2000
GroupKey[cpu, system]		

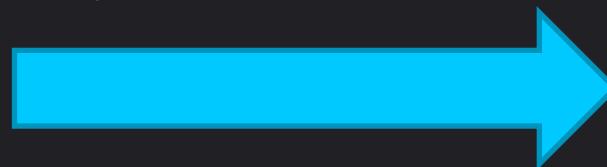
_m	_field	_value
cpu	idle	233,33
GroupKey[cpu, idle]		

NOTE: time and host columns get deleted

# Example of selector

_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
11:45	cpu	serverA	system	3000
GroupKey[cpu, system]				

| > last()



_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400
GroupKey[cpu, idle]				

_time	_m	host	_field	_value
11:45	cpu	serverA	system	3000
GroupKey[cpu, system]				

_time	_m	host	_field	_value
11:50	cpu	serverB	idle	400
GroupKey[cpu, idle]				

Let's get dirty!

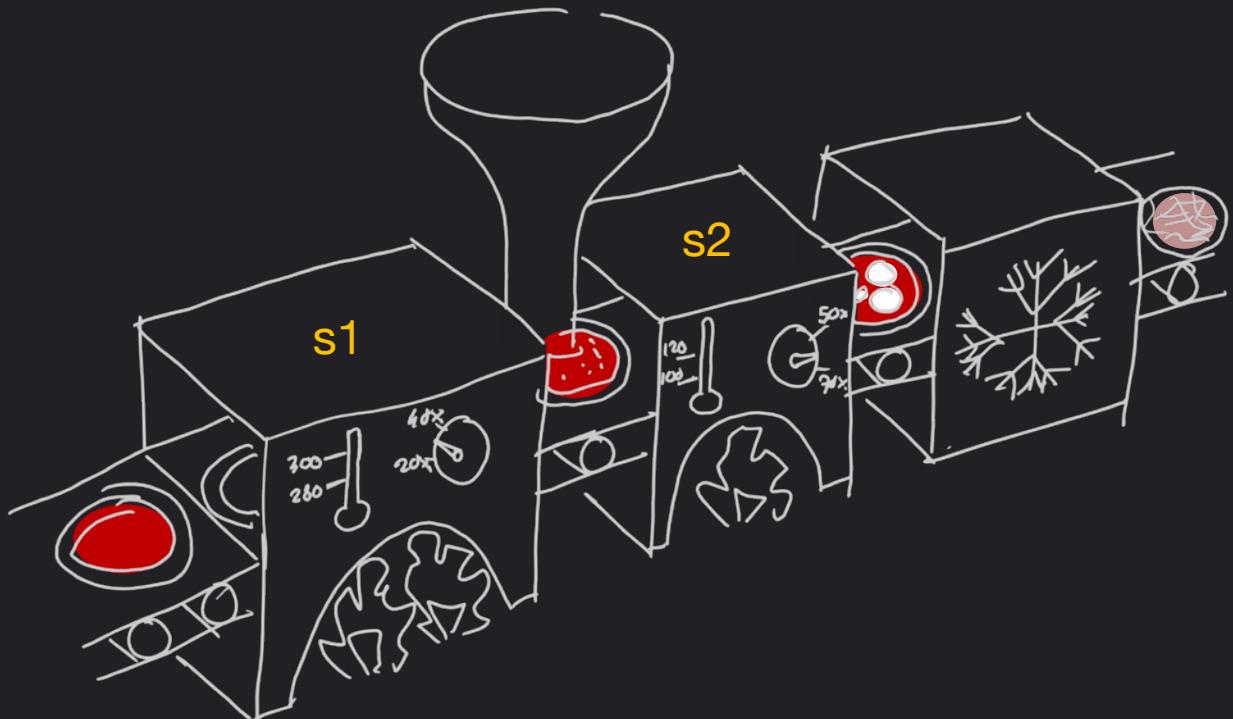


5

# Continuous Linear Pizza Oven

## Learning goals:

- Function
  - group()
  - Aggregates
  - Selectors



## Task

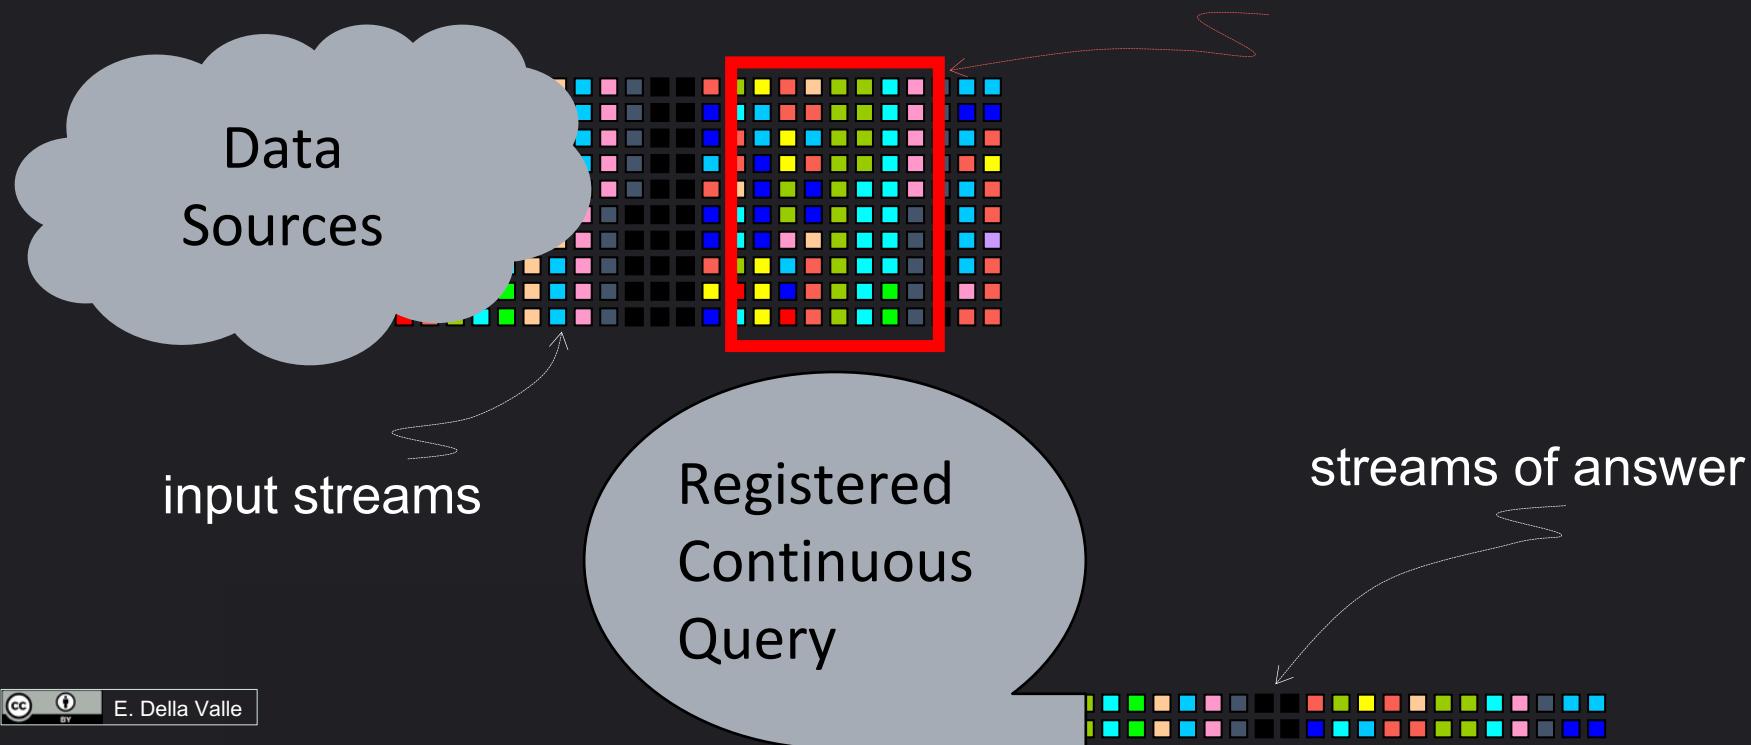
- **Grouping + Aggregator** - Extract the average temperature and the average humidity along the different stages of the linear pizza oven
- **Selector** - Extract the last humidity observation from the cooking base area

## Take home message

- **Grouping shape** tables
- **Aggregators** (e.g. mean) **compress** each table **in a row**
- **Selectors** (e.g. last) select **specific row(s)** from each table

# Exploring flux Advanced Windowing

## Recall



# Windows

- Streams are infinite in nature
- Processing can end only if the input is finite
- A **window operator** selects a finite portion of a stream
- We have seen the **time range** operator
- Flux includes a window operator that builds table **tumbling** across the results of the time range operator
  - e.g. `aggregateWindow(every: 60m, fn: mean)`

# Example of tumbling window

- Anatomy of a Flux query:

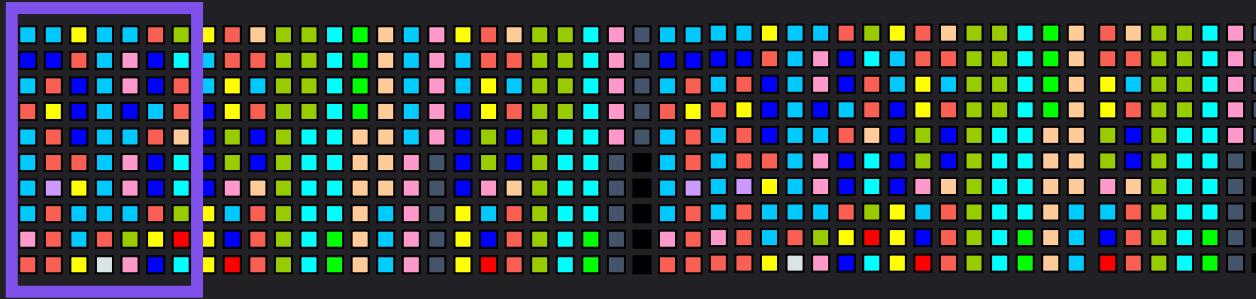
- data source                         `from(bucket:"telegraf/autogen")`
- time range                         `|> range(start:-1h)`
- data filters                         `|> filter(fn: (r) =>`
- shaping                                 `r._measurement == "cpu" and`
- processing                             `r._field == "system" and`  
   `r._value > 2000.0)`

`|> aggregateWindow(every: 60m, fn: mean)`

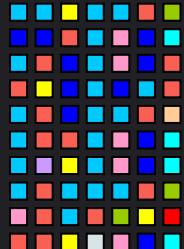
# Tumbling?



# Tumbling windows

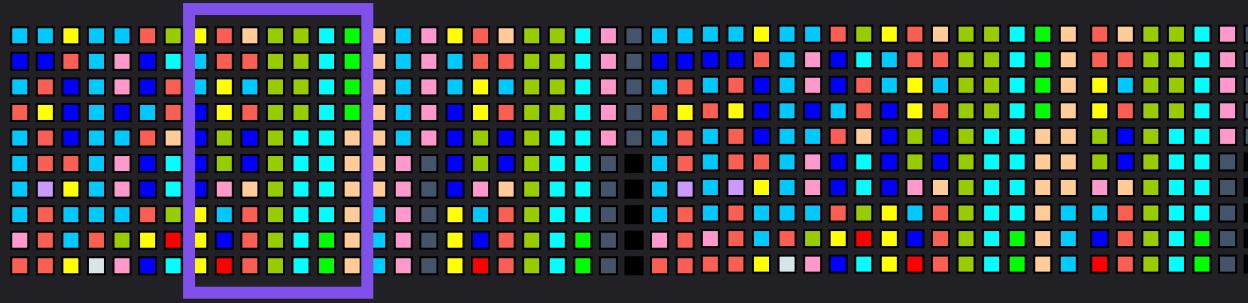


E. Della Valle

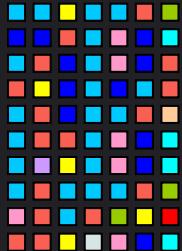


[1]

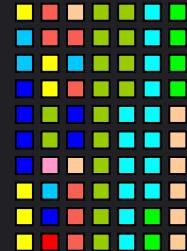
# Tumbling windows



E. Della Valle

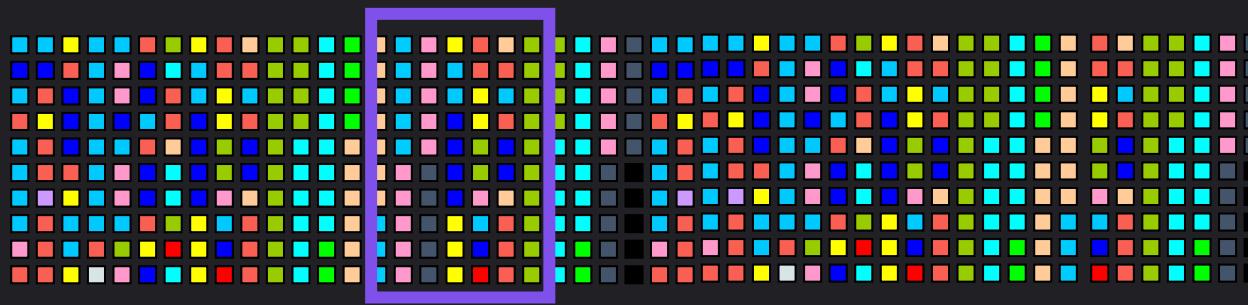


[1]



[2]

# Tumbling windows



E. Della Valle

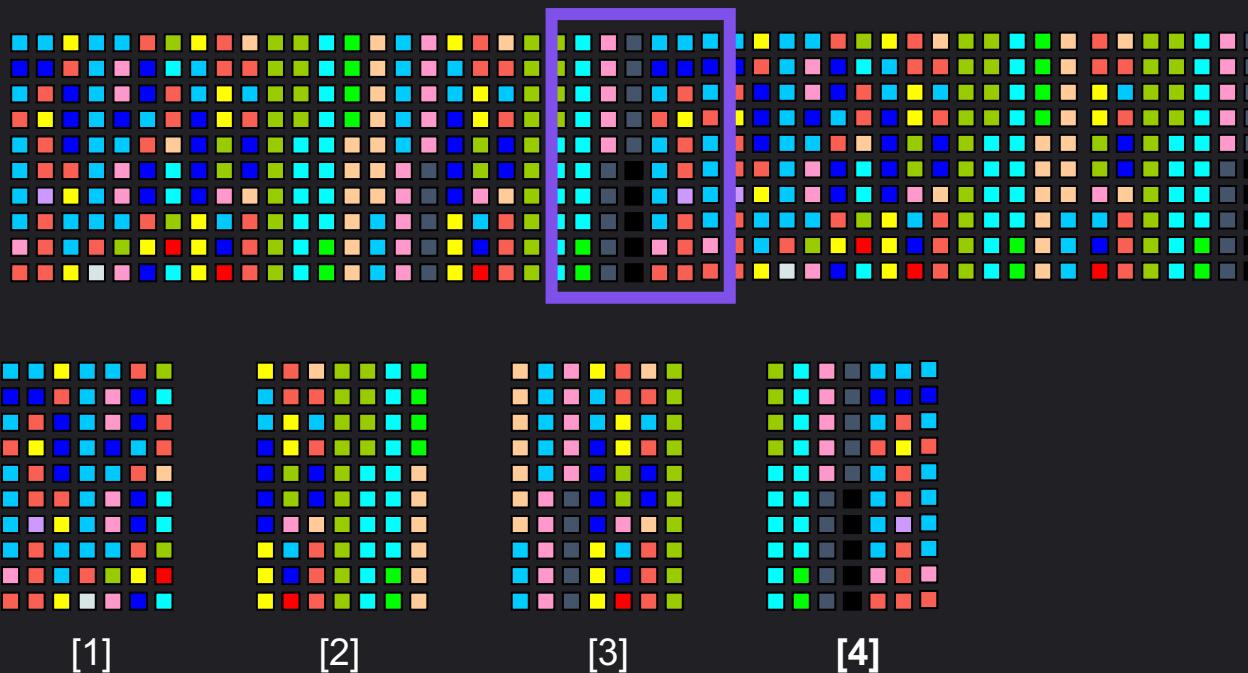


[1]

[2]

[3]

# Tumbling windows



E. Della Valle

NOTE: host column gets deleted

_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
11:45	cpu	serverA	system	3000

GroupKey[cpu, system]

_time	_m	_field	_value
11:00	cpu	system	1000
GroupKey[cpu, system]			

```
|> aggregateWindow(  
    every: 60m,  
    fn: mean)
```

_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400

GroupKey[cpu, idle]

_time	_m	_field	_value
11:00	cpu	idle	100
GroupKey[cpu, idle]			

NOTE: The start time is not included, the stop time is

NOTE: host column gets deleted

_time	_m	host	_field	_value
10:45	cpu	serverA	system	1000
11:00	cpu	serverB	system	2000
11:30	cpu	serverB	system	2000
11:45	cpu	serverA	system	3000

GroupKey[cpu, system]

```
|> aggregateWindow(  
    every: 60m,  
    fn: mean)
```

_time	_m	_field	_value
11:00	cpu	system	1000
GroupKey[cpu, system]			

_time	_m	host	_field	_value
10:35	cpu	serverA	idle	100
11:15	cpu	serverB	idle	200
11:50	cpu	serverB	idle	400

GroupKey[cpu, idle]

NOTE: The start time is not included, the stop time is

_time	_m	_field	_value
12:00	cpu	system	2333.3
GroupKey[cpu, system]			

_time	_m	_field	_value
11:00	cpu	idle	100
GroupKey[cpu, idle]			

_time	_m	_field	_value
12:00	cpu	idle	300
GroupKey[cpu, idle]			

## Typical use case: downsampling data



## Typical use case: downsampling data (cont.)



`aggregateWindow(every: 5m, fn: mean)`



# Let's get dirty!

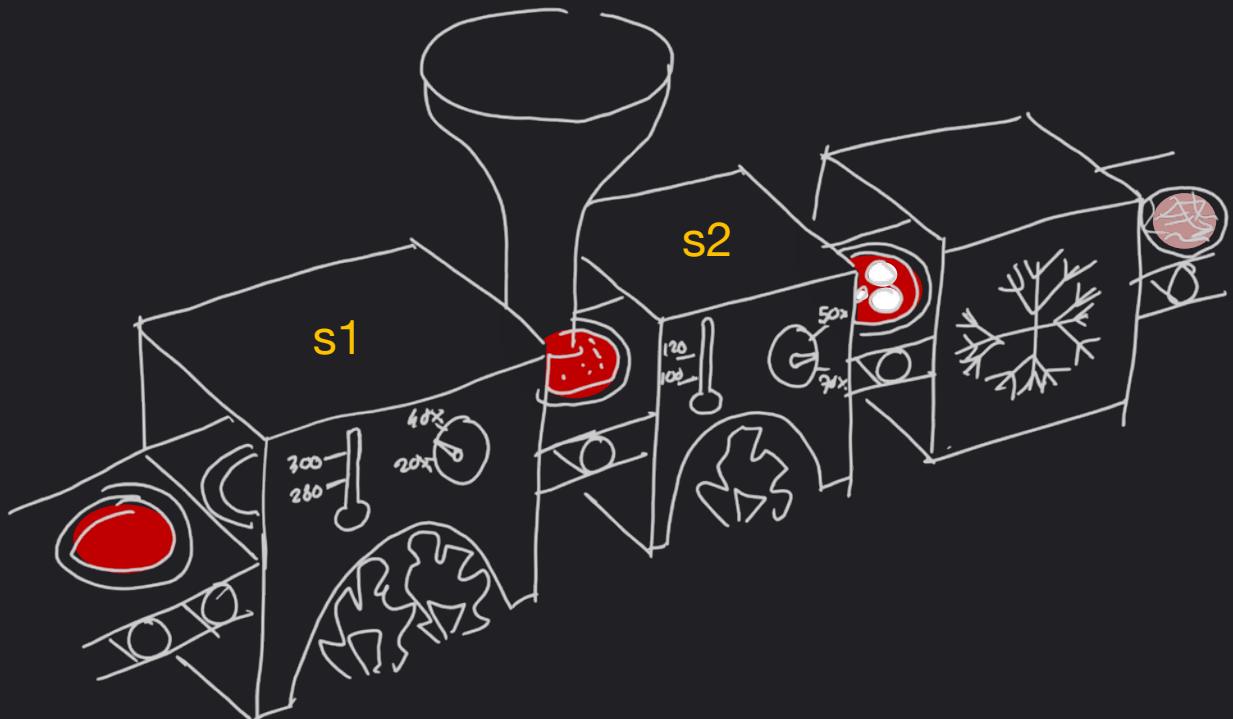


6

# Continuous Linear Pizza Oven

Learning goals:

- aggregateWindow



## Task

- Extract the moving average of the temperatures observed in the cooking base area over a window of 2 minutes
- Extract the moving average of the temperatures observed by S2 over a window of 3 minutes

# Quiz

- True or false
  - you can omit the range clause – ?
  - filtering can change the number of tables in the result – ?
  - shaping always changes the number of tables in the result – ?
  - processing can change the number of tables in the result – ?
  - aggregateWindow function is just a process step – ?
  - aggregateWindow function eases the downsampling task – ?
  - the parameter createEmpty of the aggregateWindow function can change the number of tables in the result – ?

## Quiz answer

- True or false
  - you can omit the range clause – F
  - filtering can change the number of tables in the result – T (empty tables)
  - shaping always changes the number of tables in the result – F
  - processing can change the number of tables in the result – F
  - aggregateWindow function is just a process step – F
  - aggregateWindow function eases the downsampling task – T
  - the parameter createEmpty of the aggregateWindow function can change the number of tables in the result – T



*influx*/days



## Basic Data Analysis

---

**Emanuele Della Valle** Prof. @ Politecnico di Milano & Partner @ Quantia Consulting

**Marco Balduini** Founder & CEO @ Quantia Consulting

**Riccardo Tommasini** Prof. @ University of Tartu