



influx/days

Data Analysis

Emanuele Della Valle

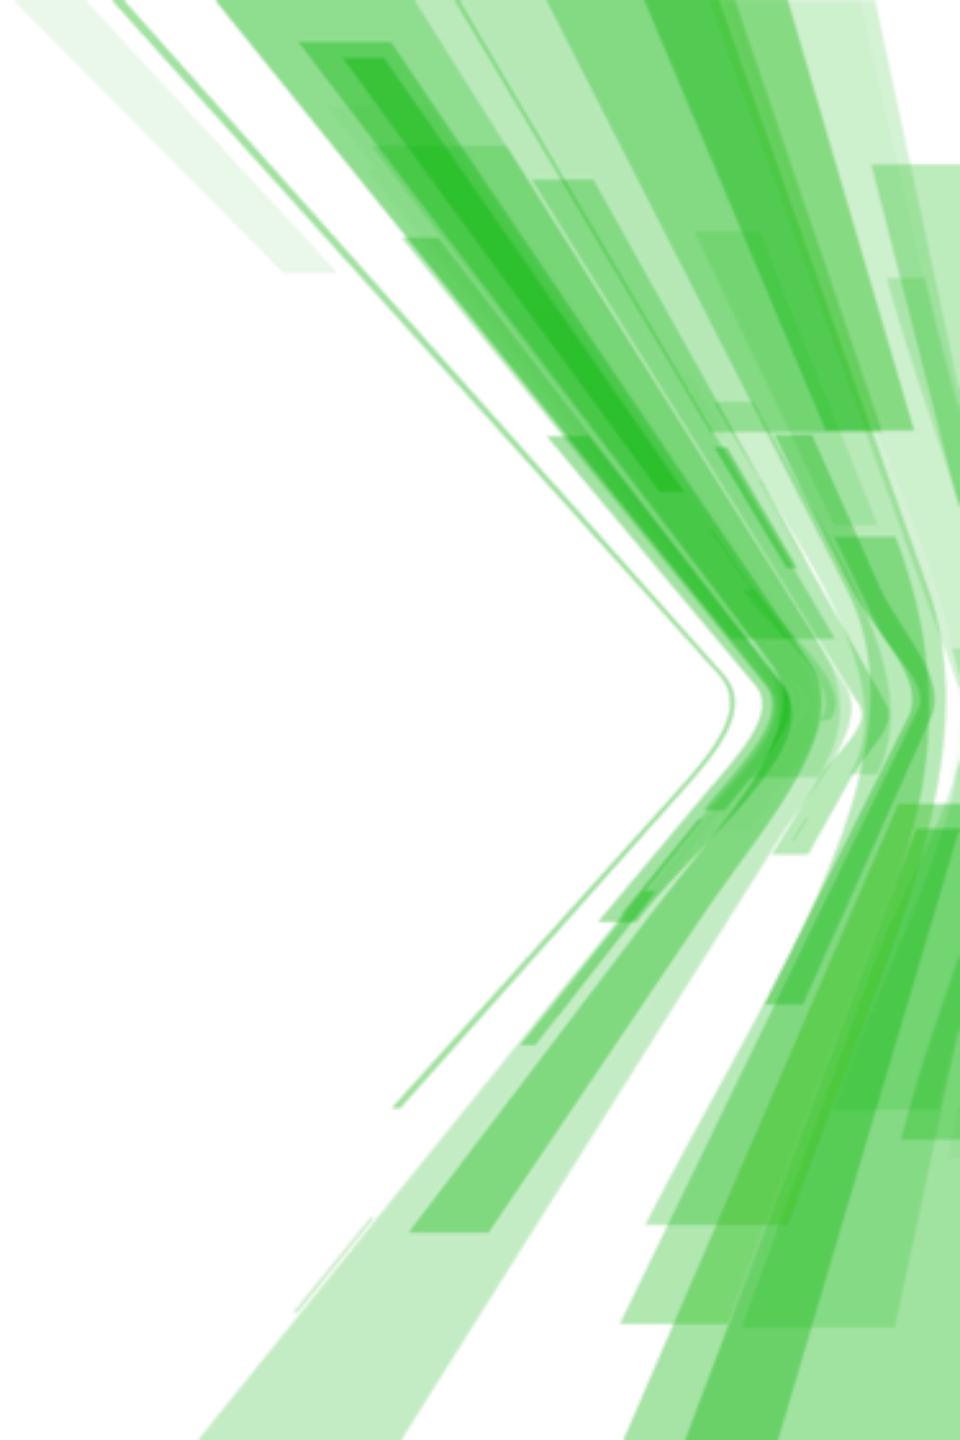
Prof. @ Politecnico di Milano & Partner
@ Quantia Consulting

Marco Balduini

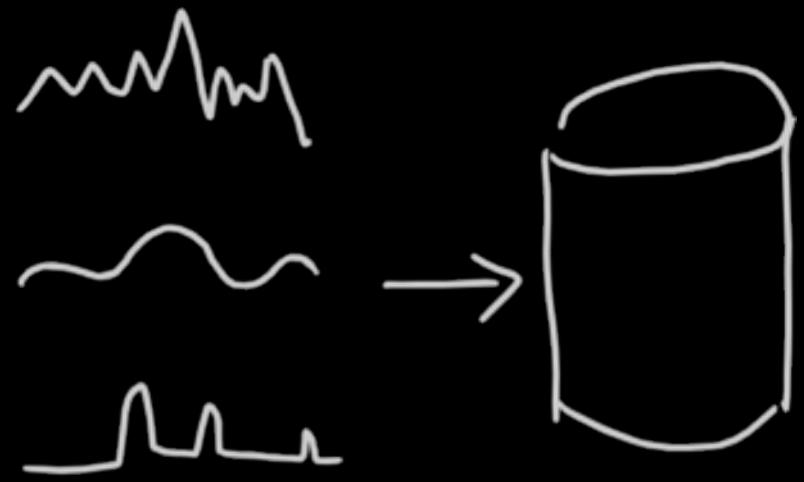
Founder & CEO @ Quantia Consulting



Introduction



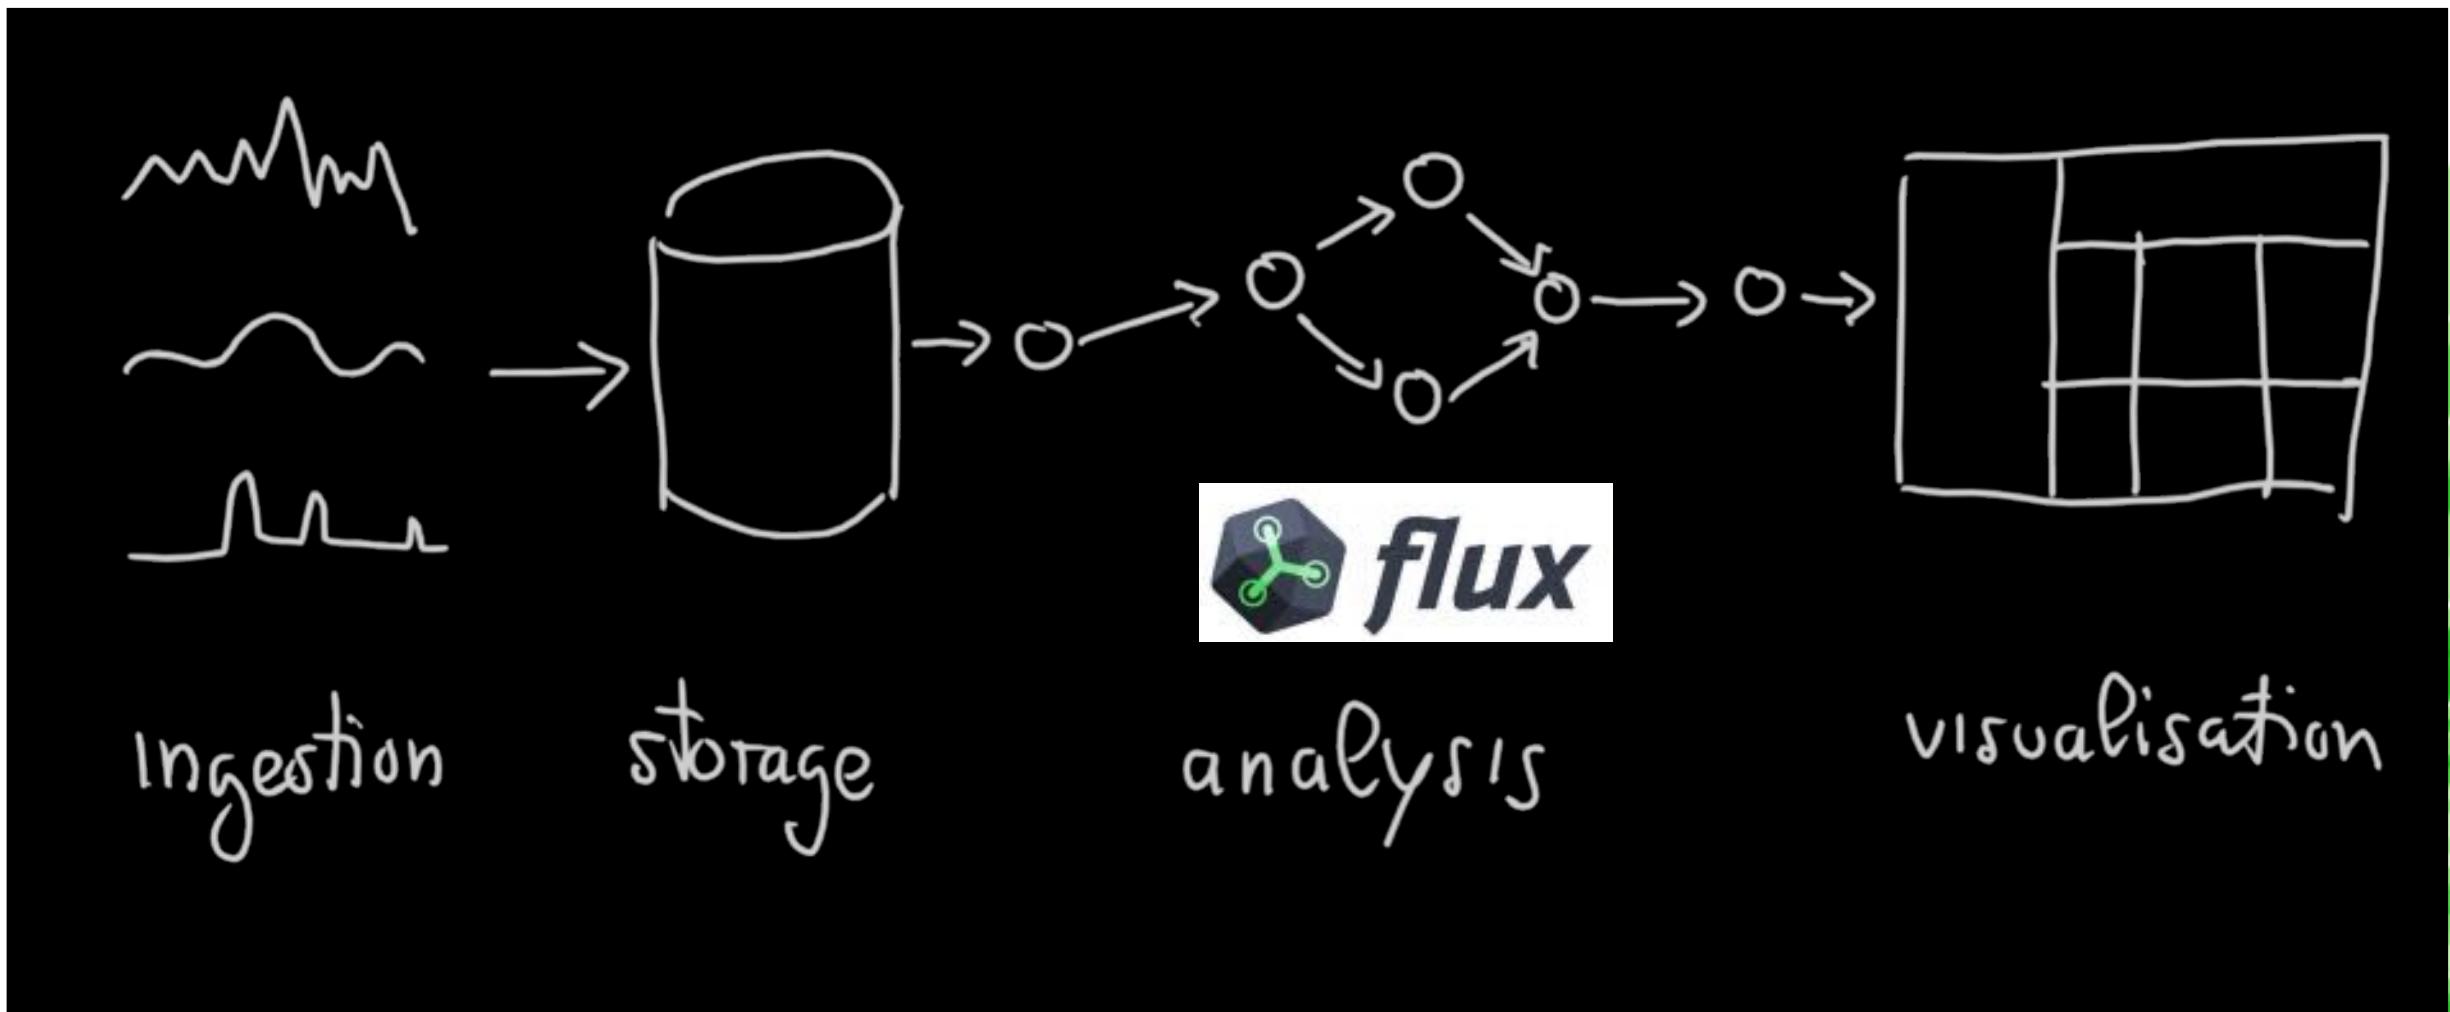
Data Lifecycle



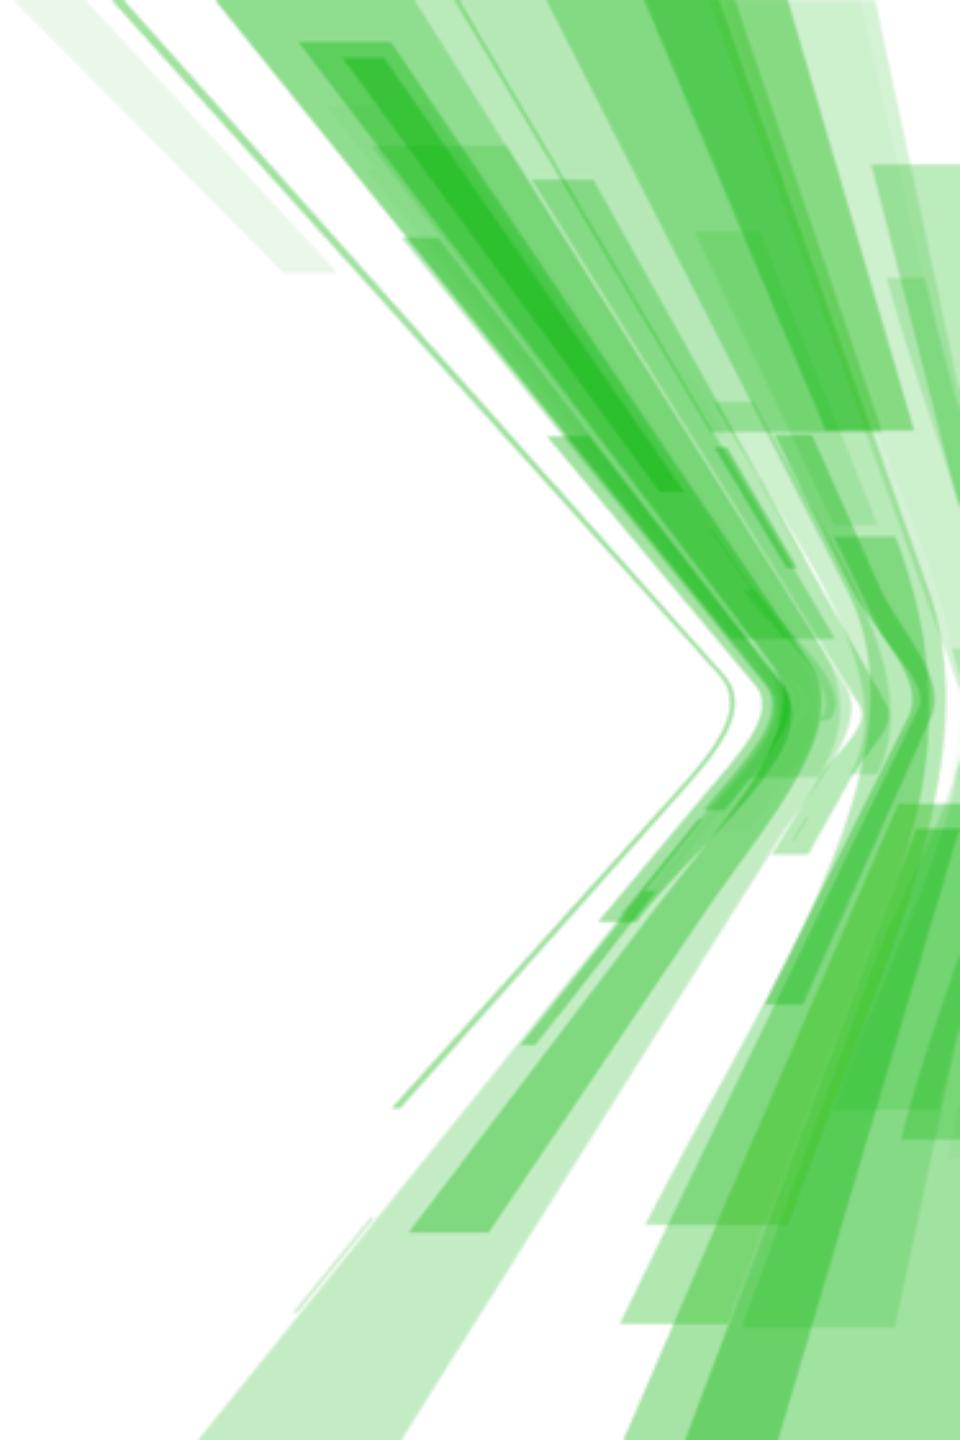
ingestion

storage

Data Lifecycle



Flux



What Is Flux?

Flux is a functional data scripting and query language

Written to be:

- Useable: easy to learn
- Readable: developers read more code than we write
- Composable: developers can build onto the language
- Testable: queries are code
- Contributable: open source contributions matter
- Shareable: developers read more code than we write

What is Flux? (cont.)

- Open Source (MIT)
- Data-first mentality
- Streaming Data
- Planner-optimized
- Compatible with multiple sources and sinks

Flux vs. InfluxQL

Flux is an alternative to InfluxQL with additional functionality like:

- Joins
- Math across measurements
- Sort by any column
- Pivot
- Histograms
- Covariance

Reference: <https://v2.docs.influxdata.com/v2.0/query-data/>

Basic Flux Language Elements

Every Flux query needs:

- data source
- time range
- data filters

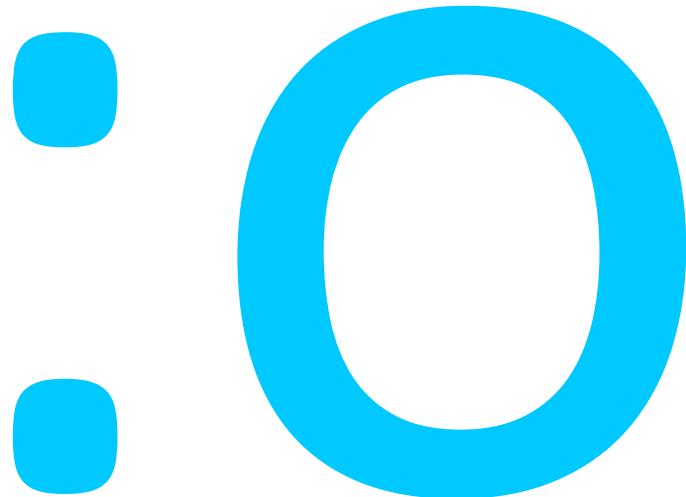
Pipe-forward operator:
chain operations together

```
from(bucket:"telegraf/autogen")  
|> range(start:-1h)  
|> filter(fn: (r) =>  
  r._measurement == "cpu" and  
  r._field == "system" and  
  r._value > 2000.0 )
```

Anonymous function

Flux Data Model

An infinite stream...
of finite tables...
identified by the GroupKey



From bucket to tables

_time	host	_field	_value
1492...0	serverA	system	2342.2
1492...	serverA	system	477.4
GroupKey[host=serverA,_field=system]			

_time	host	_field	_value
1492...1	serverB	system	2.779
1492...	serverB	system	1274.5
GroupKey[host=serverB,_field=system]			

_time	host	_field	_value
1492...1	serverB	idle	1789.4
1492...	serverB	idle	2475.9
GroupKey[host=serverB,_field=idle]			



from(bucket:"telegraf/autogen")

|> range(start:-1s)

_time	host	idle	system
1492...0	serverA	1.667	2342.2
...
1492...1	serverB	1789.4	2.779

Query Process

For Each Table:

1. Convert incoming rows into 0 or more outgoing rows
2. Map outgoing rows to tables

_time	host	_field	_value
1492...0	serverA	system	2342.2
1492...	serverA	system	477.4

GroupKey[host=serverA,_field=system]

_time	host	_field	_value
1492...1	serverB	system	2.779
1492...	serverB	system	1274.5

GroupKey[host=serverB,_field=system]

_time	host	_field	_value
1492...1	serverB	idle	1789.4
1492...	serverB	idle	2475.9

GroupKey[host=serverB,_field=idle]

```
|> filter(fn: (r) =>  
  r._measurement == "cpu" and  
  r._field == "system" and  
  r._value > 2000.0)
```



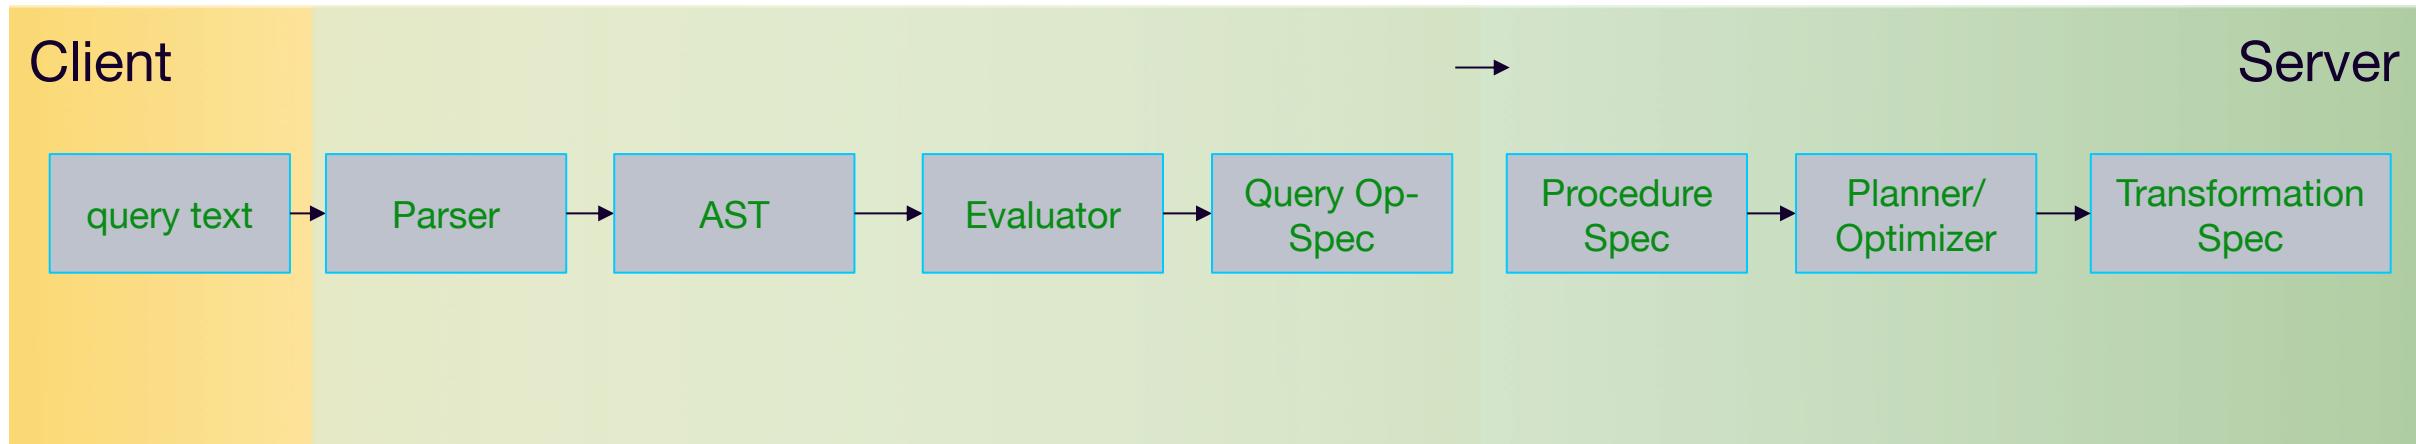
_time	host	_field	_value
1492...0	serverA	system	2342.2

GroupKey[host=serverA,_field=system]

_time	host	_field	_value
1492...1	serverB	system	2.779

GroupKey[host=serverB,_field=system]

The Life of a Flux Query



Wrap up of Key terms and terminology

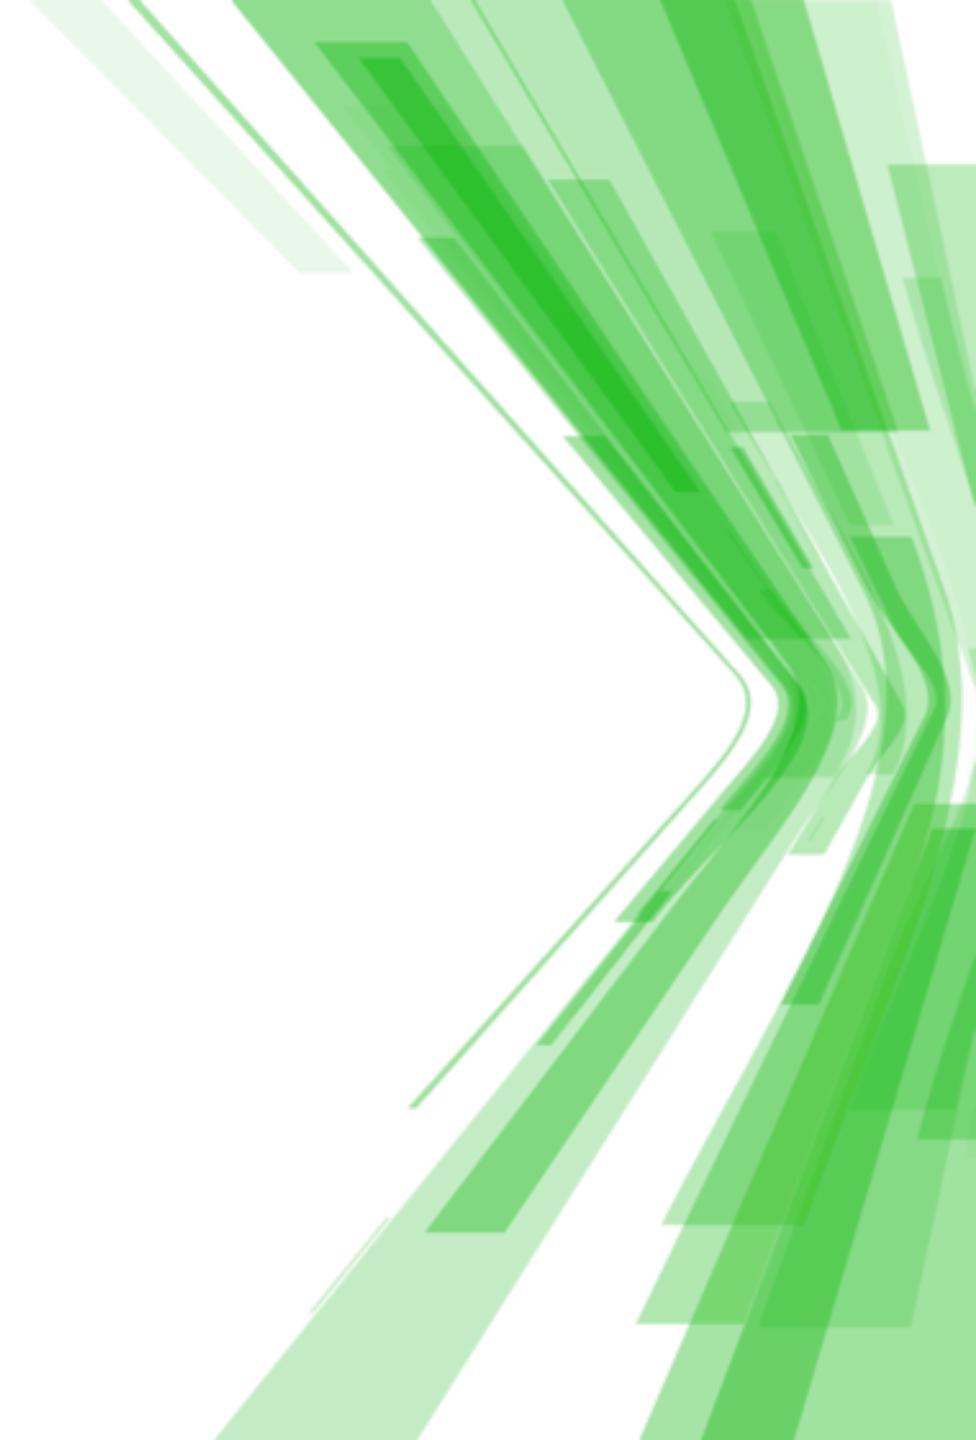
Bucket: Named data source with retention policy

Pipe-forward operator `|>` : chain operations (e.g., time range and data filters) together

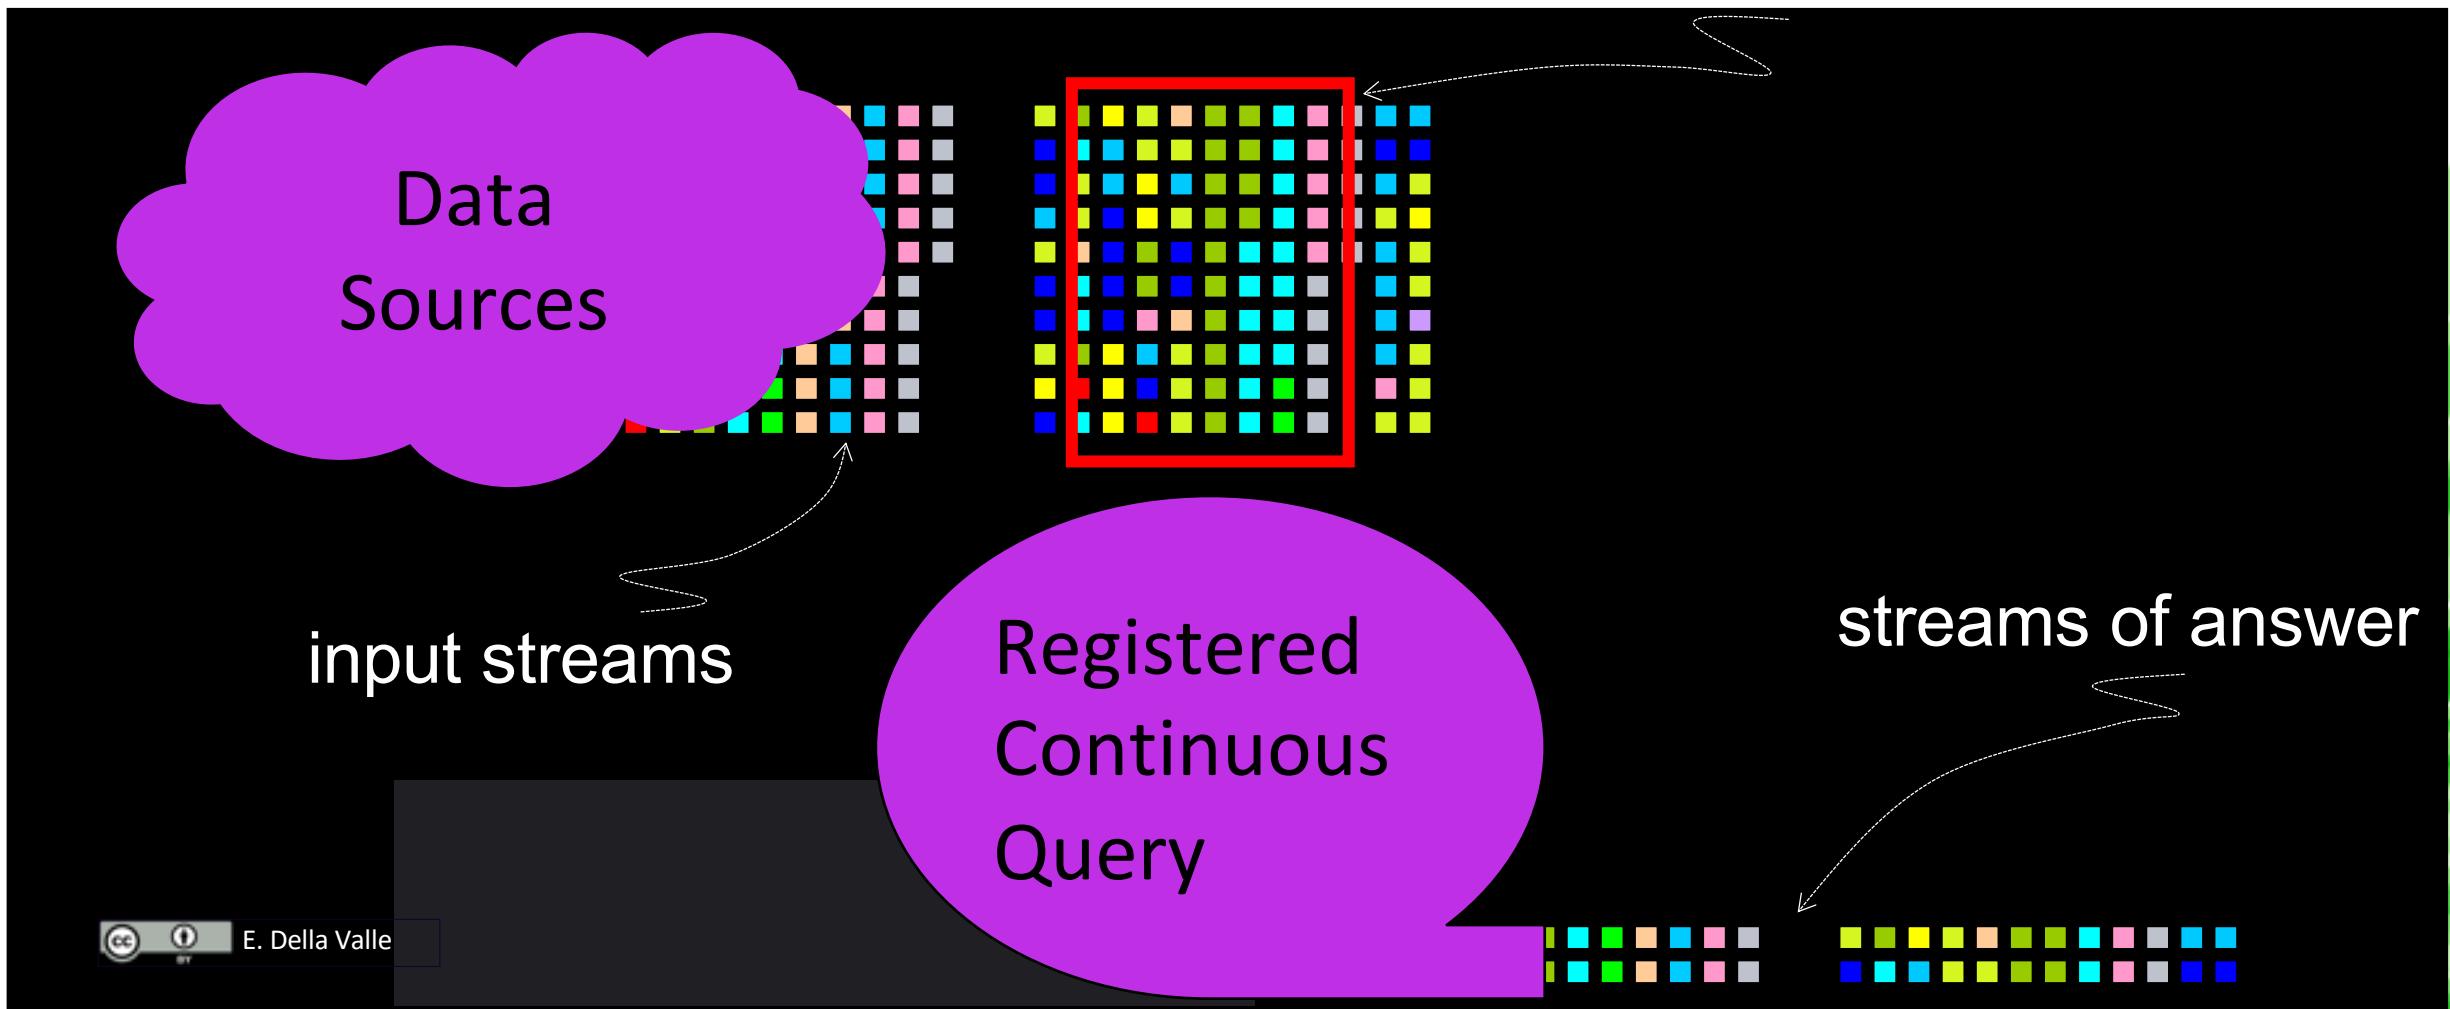
Table: data is returned in annotated CSVs, represented as tables

Group Key: list of columns for which every row in the table has the same value

Exploring Flux Basics



Key concepts



Windows

Streams are infinite in nature

Processing can end only if the input is finite

A **window operator** selects a finite portion of a stream

- **Time range** is a window operator whose input is a time interval specified either stating an **absolute** start and end time

```
range(start: 2019-10-01T00:00:00Z,  
      stop: 2019-10-01T00:05:00Z)
```

or stating a **relative** time period relative

```
range(start:-2h, stop: -1h)
```

```
range(start:-1h)
```

Let's get dirty!



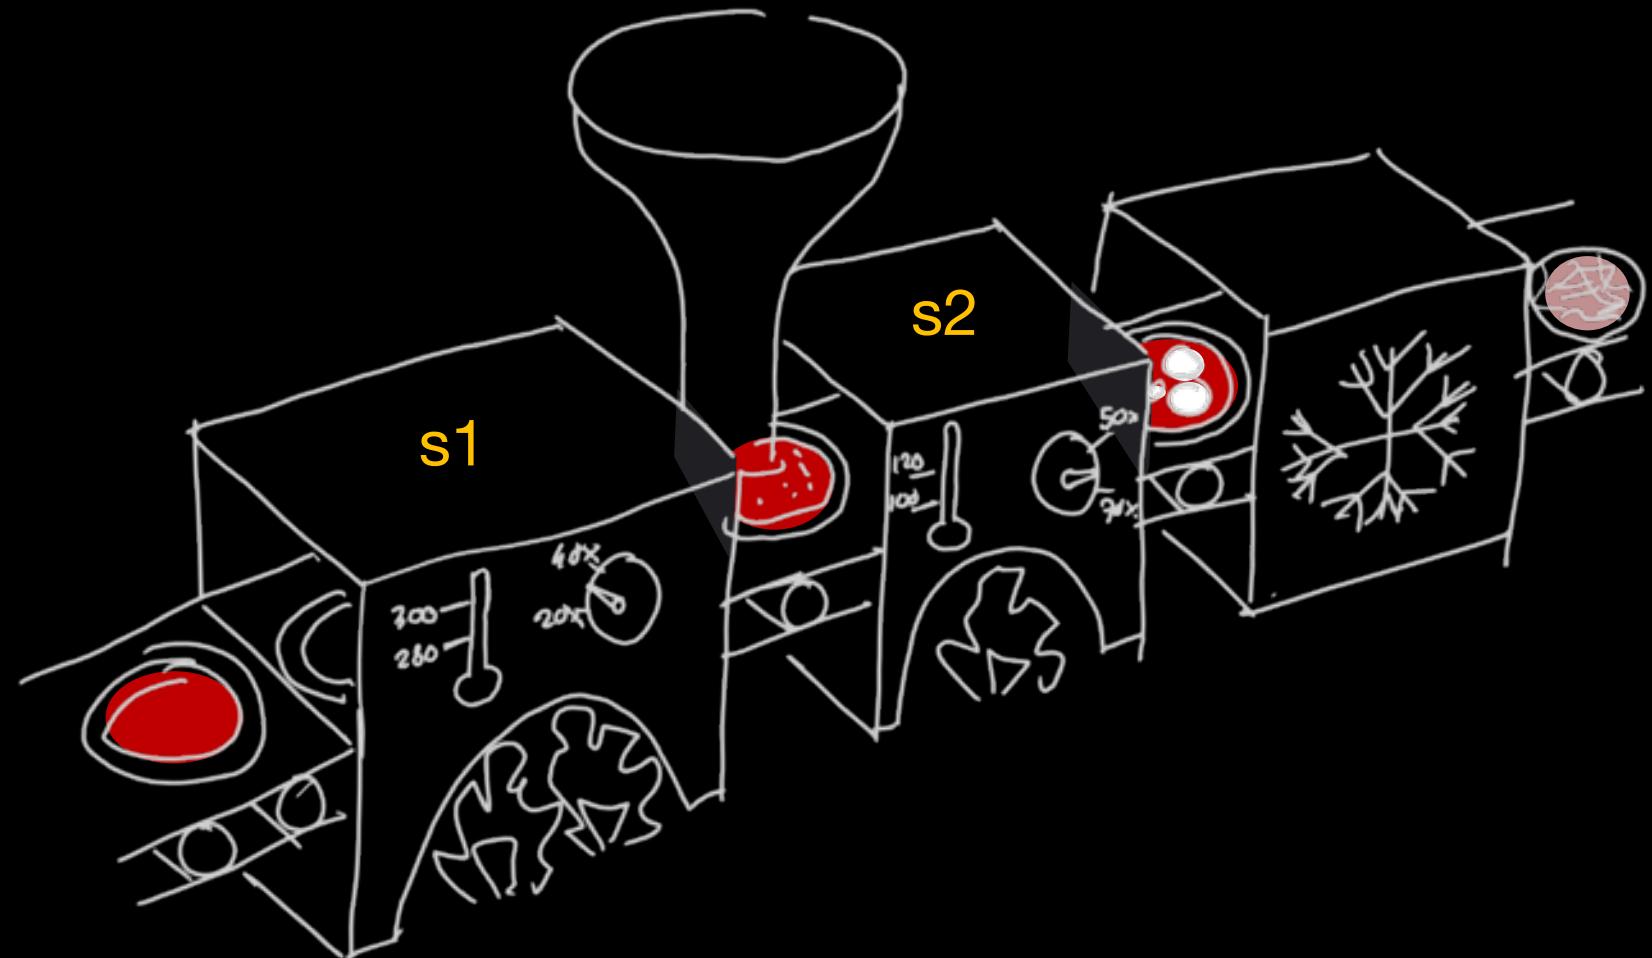
Continuous Linear Pizza Oven

Learning goals:

Data source

Window

Tables



Task

Extract all the measurements in a given range

Absolute: from 2019-10-01T00:00:00Z to 2019-10-01T00:10:00Z

Relative: in the last 24 hours

Take home message

The **range()** function **extracts** from the named bucket all the possible **tables in the given time range**

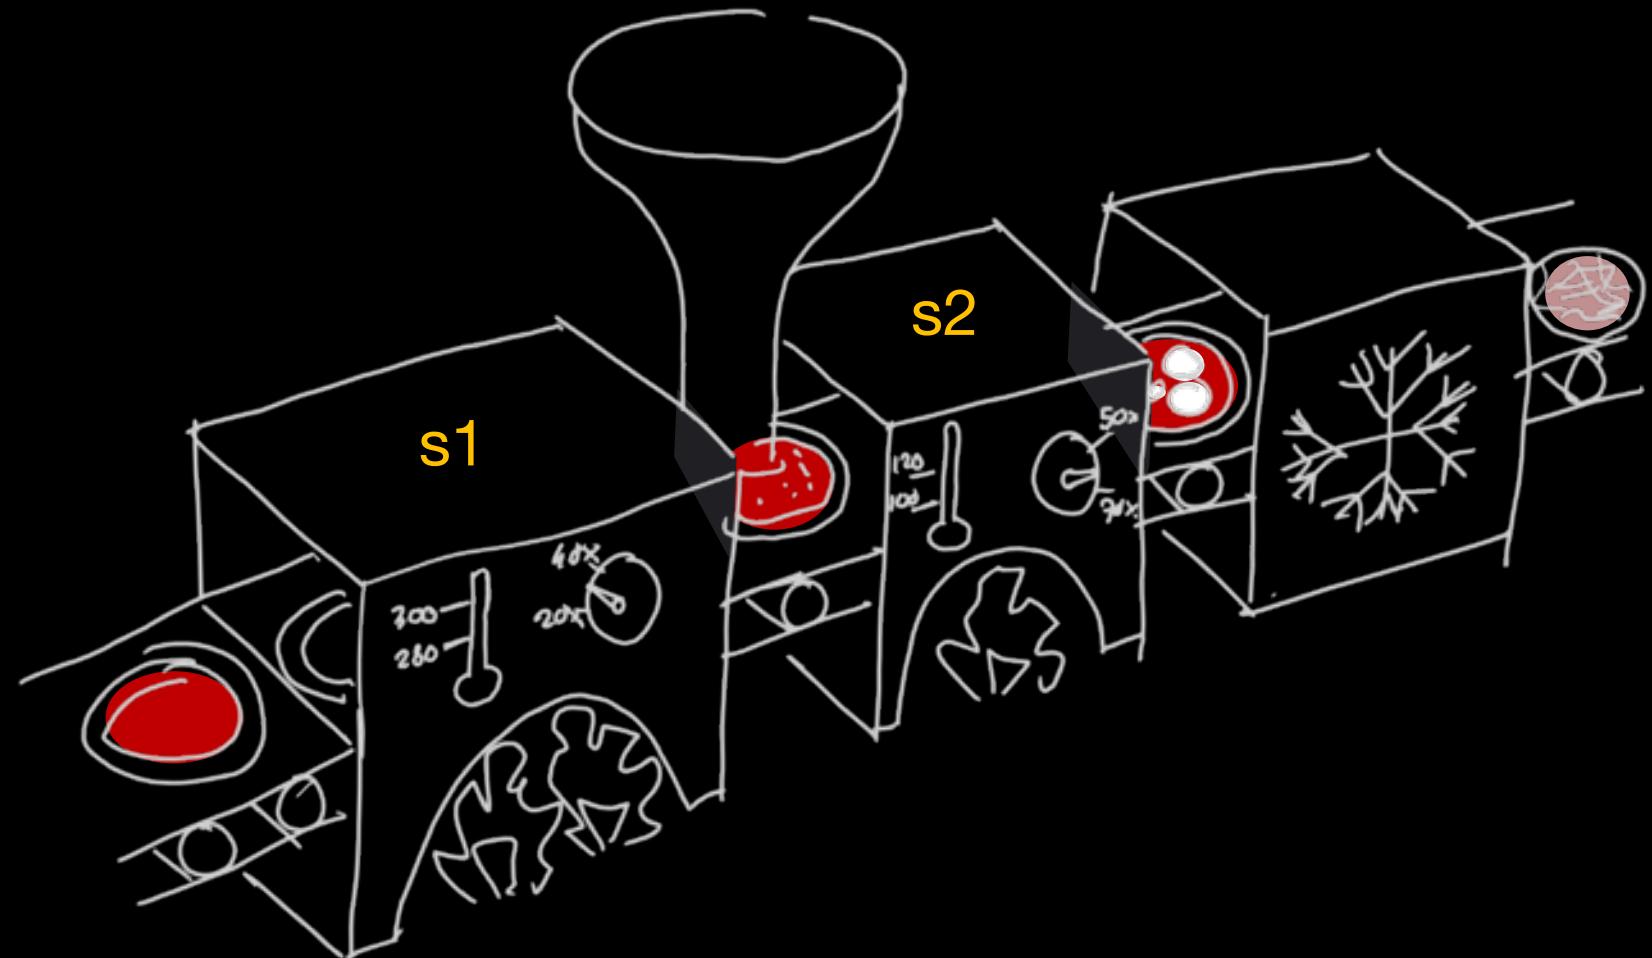
Let's get dirty!



Continuous Linear Pizza Oven

Learning goals:

Filter by tag



Task

Extract the temperature data from the cooking base area (sensor S1)

Take home message

The **filter()** function **applied to tags selects** only the **table(s)** of your interest from all the possible tables

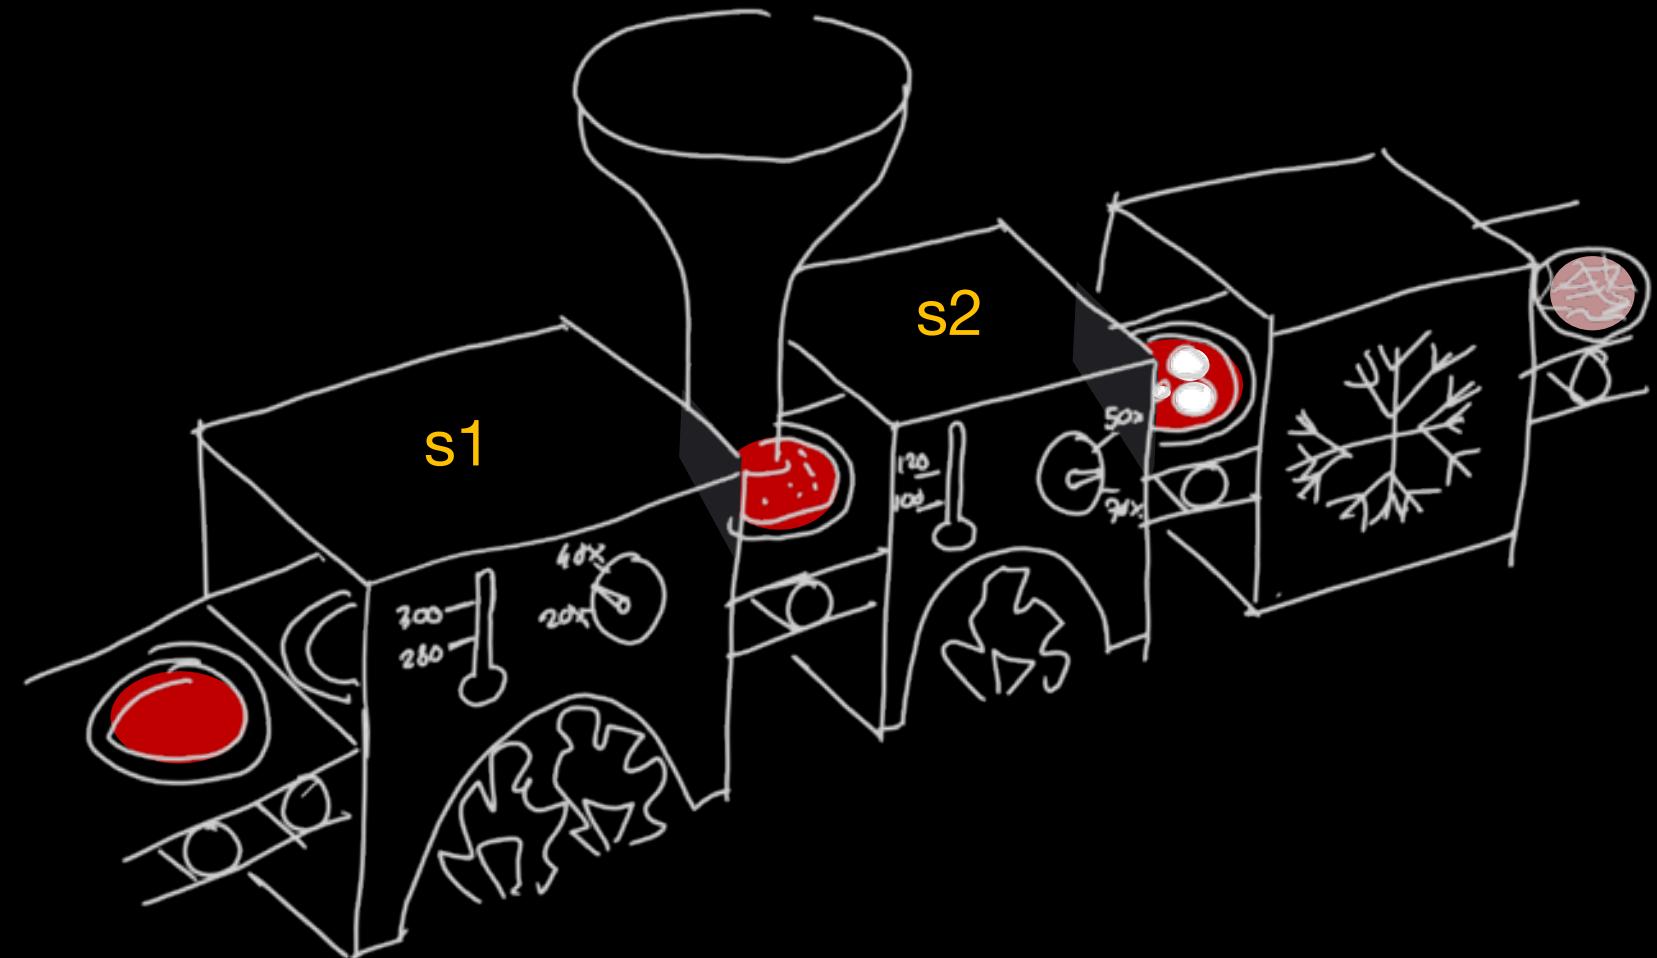
Let's get dirty!



Continuous Linear Pizza Oven

Learning goals:

Filter by value



Task

Extract the measurements from the cooking base area (sensor S1) with a temperature under 300°

Take home message

The **filter()** function **applied to values selects** only the **row(s)** of your interest from all the possible tables

Let's get dirty!

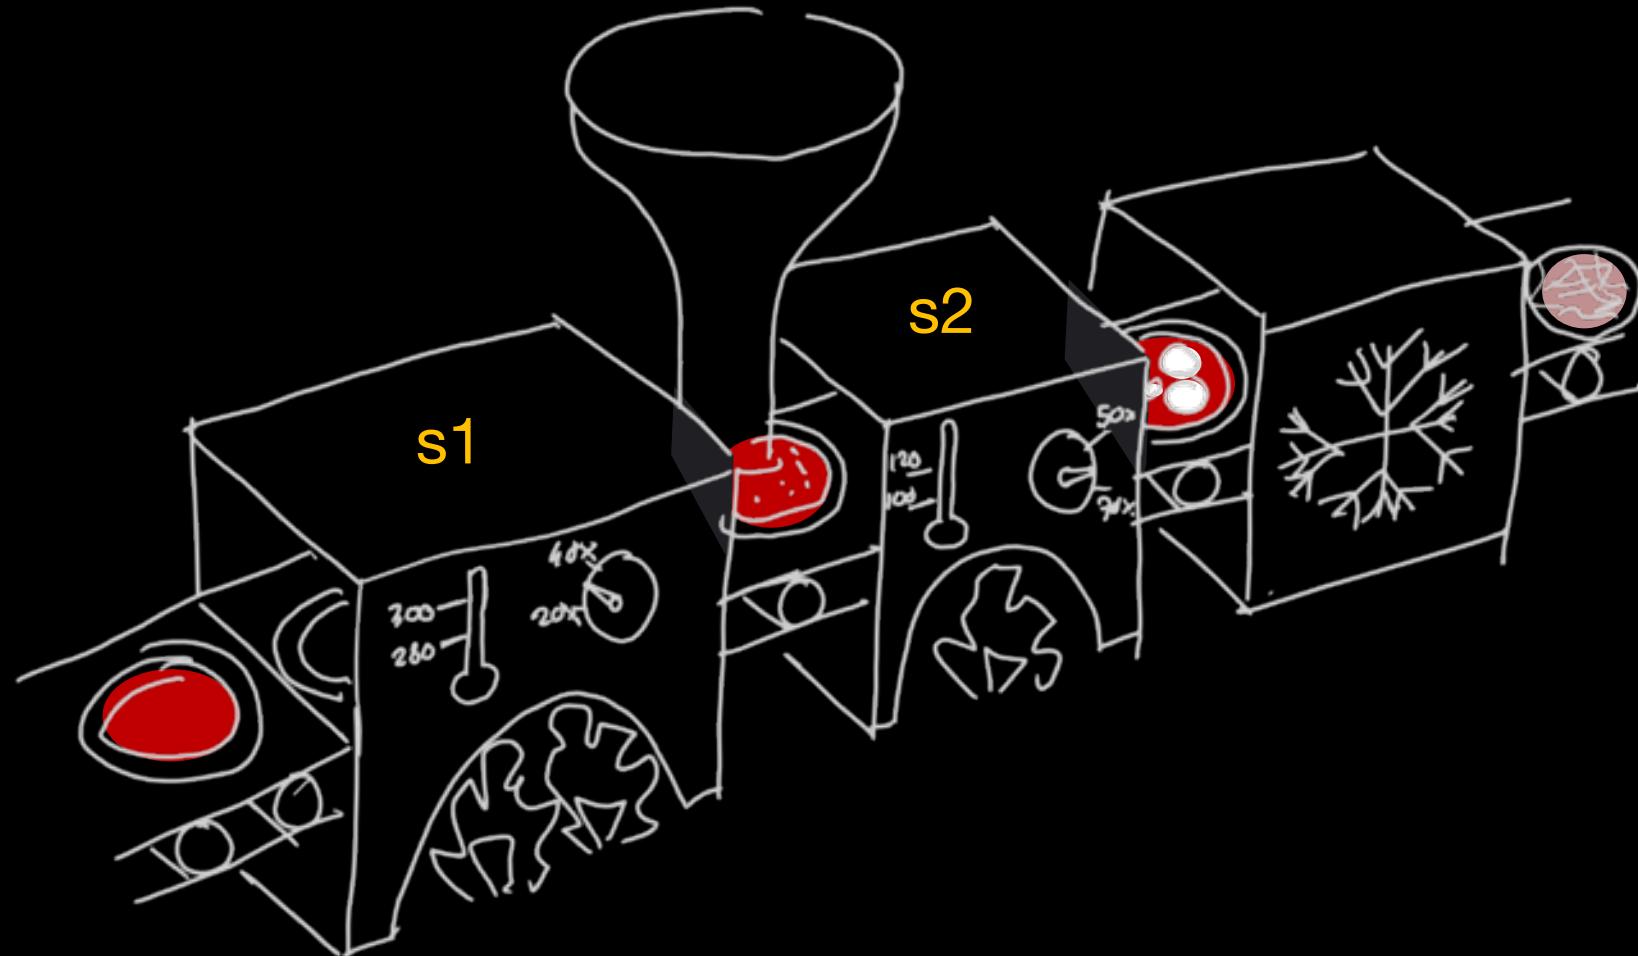


Continuous Linear Pizza Oven

Learning goals:

Function

- Aggregates
- Selectors



Task

Aggregator - Extract the average temperature in the cooking base area

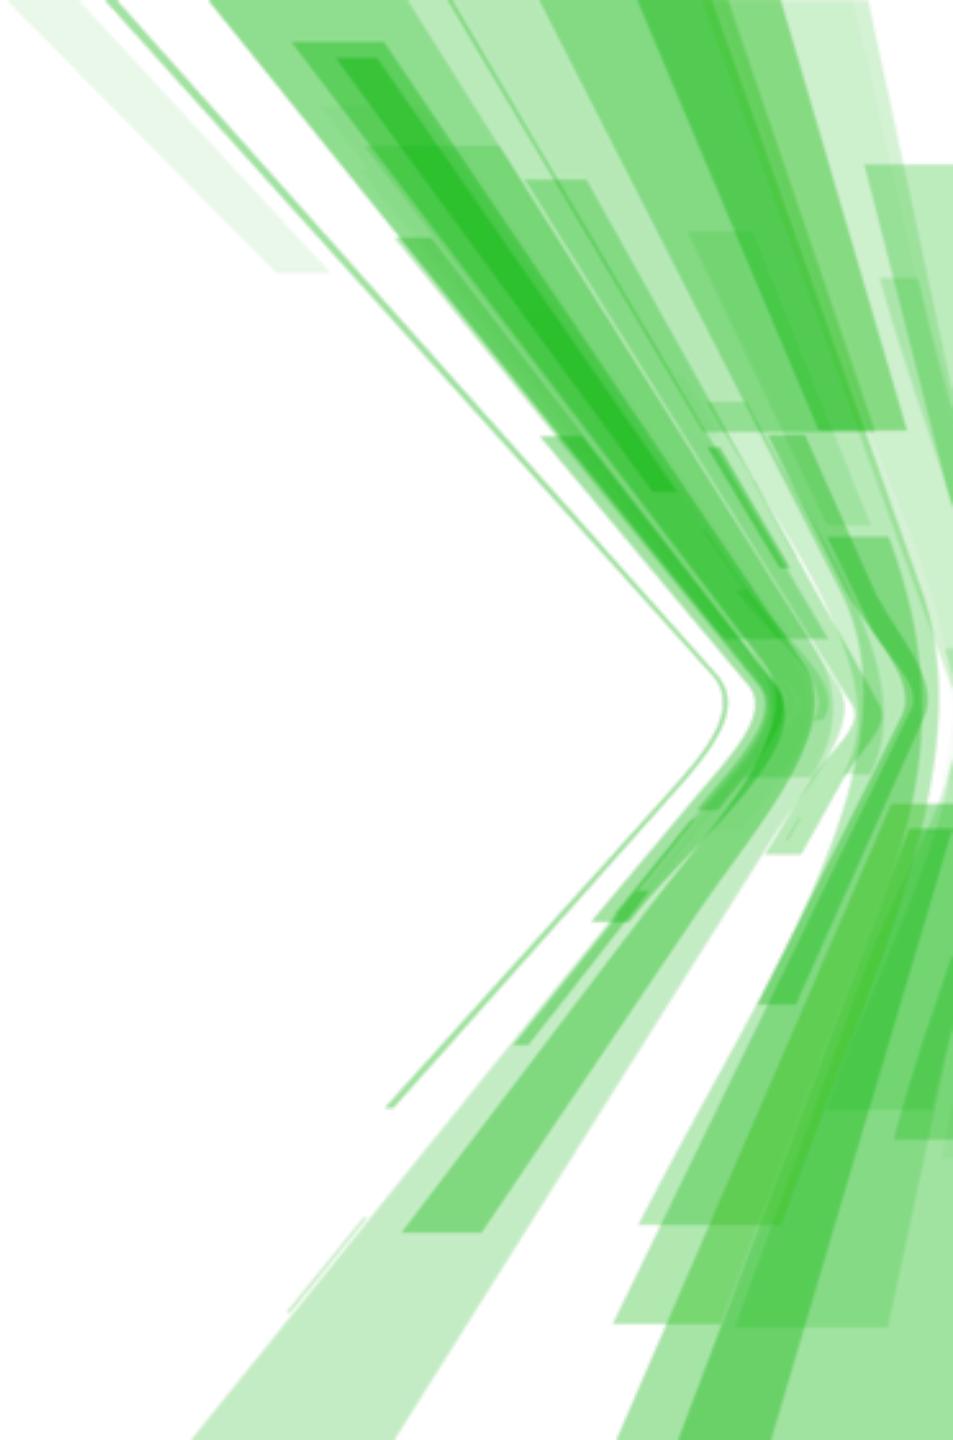
Selector - Extract the last humidity measurements from the cooking base area

Take home message

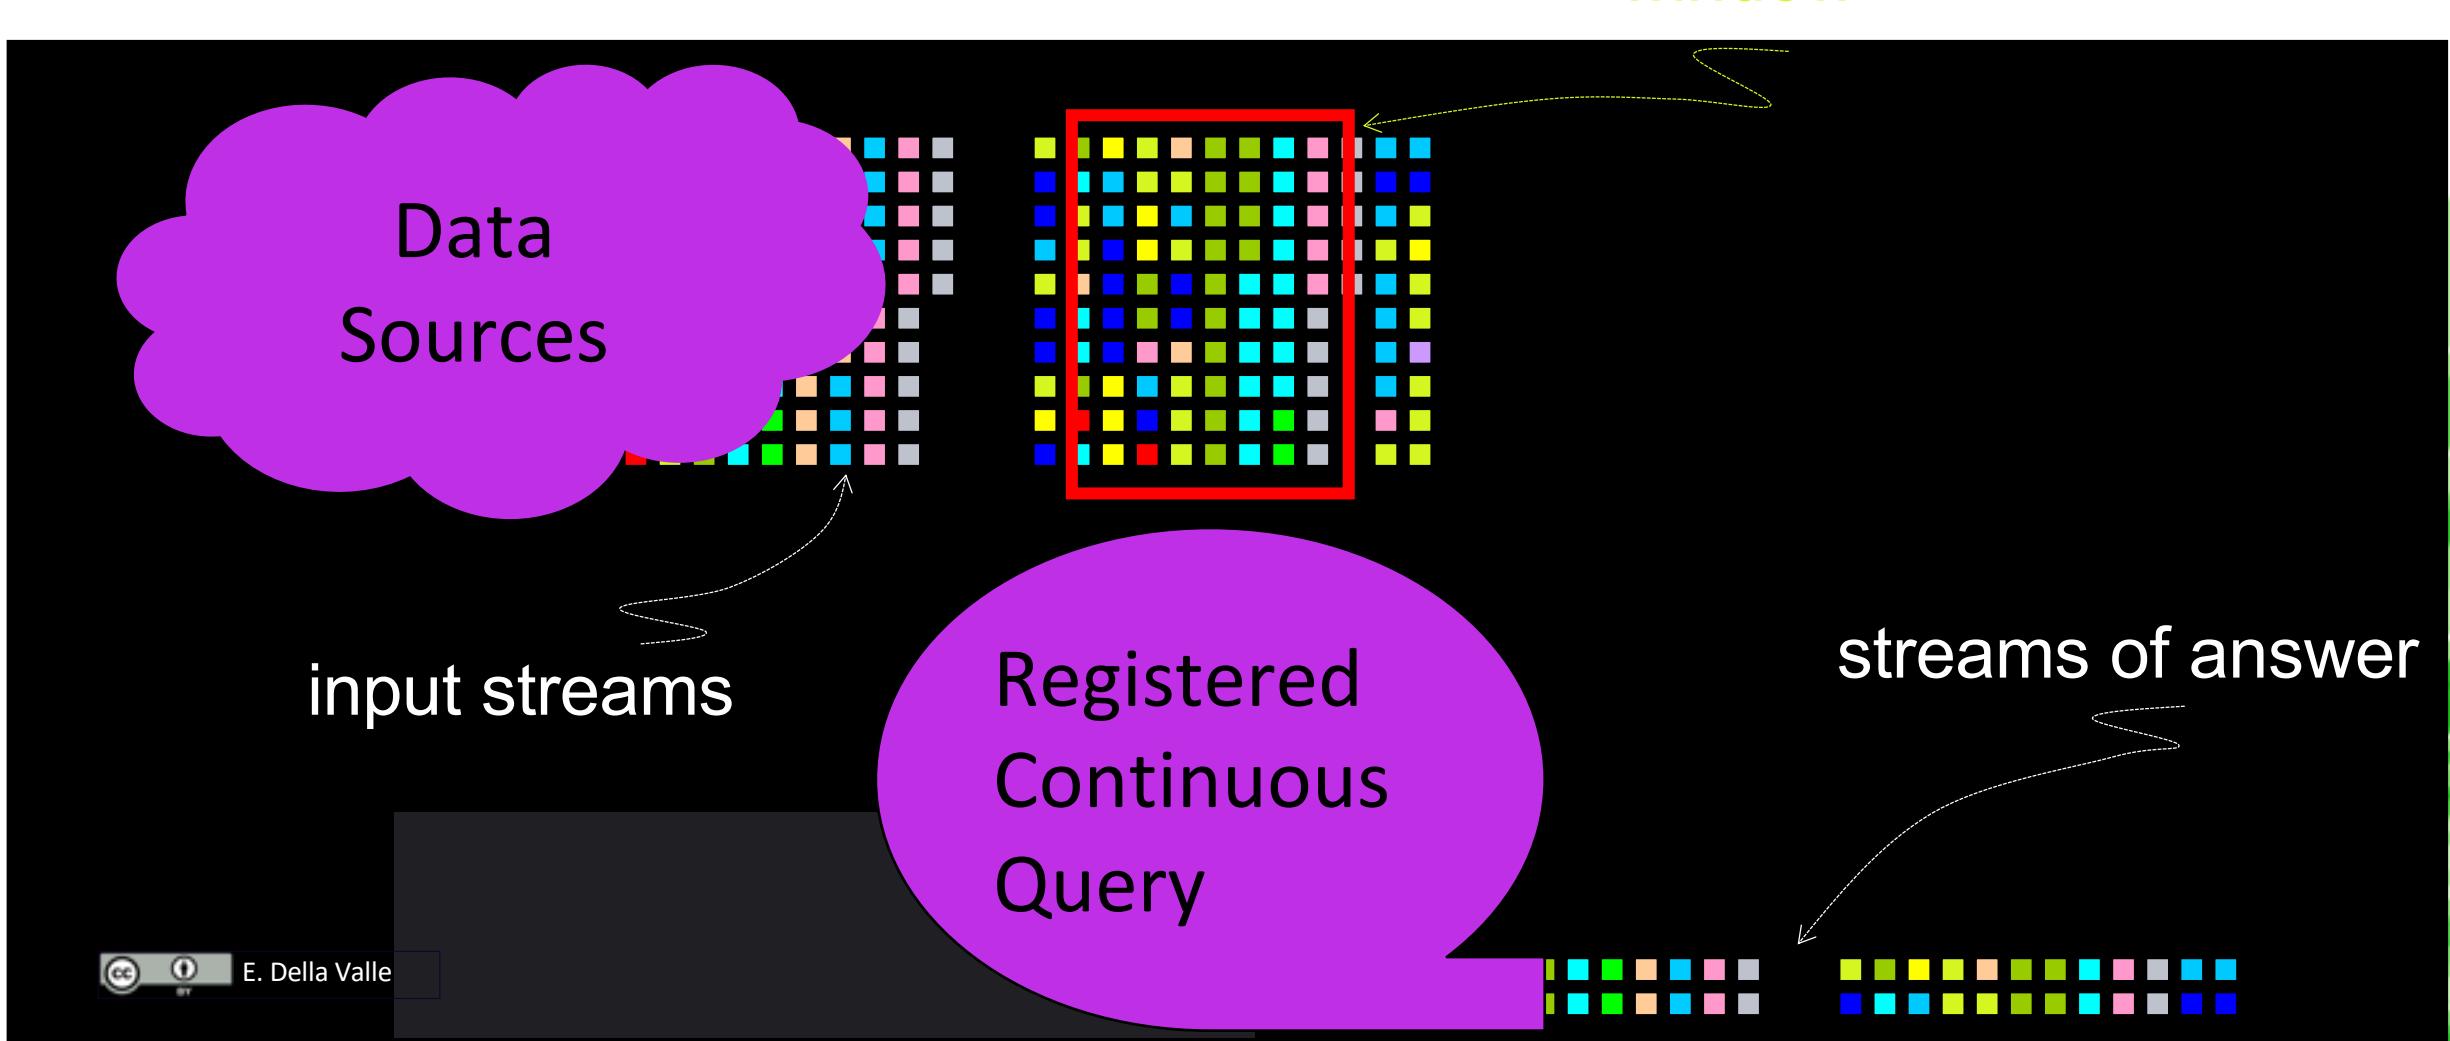
Aggregators (e.g. mean) **compress** each table **in a row**

Selectors (e.g. last) select **specific row(s)** from each table

Exploring Flux Windowing



Recall



Windows

Streams are infinite in nature

Processing can end only if the input is finite

A **window operator** selects a finite portion of a stream

We have seen the **time** range operator

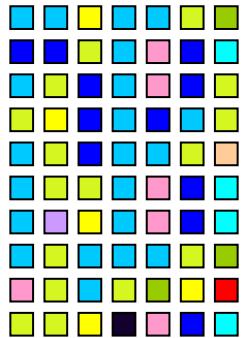
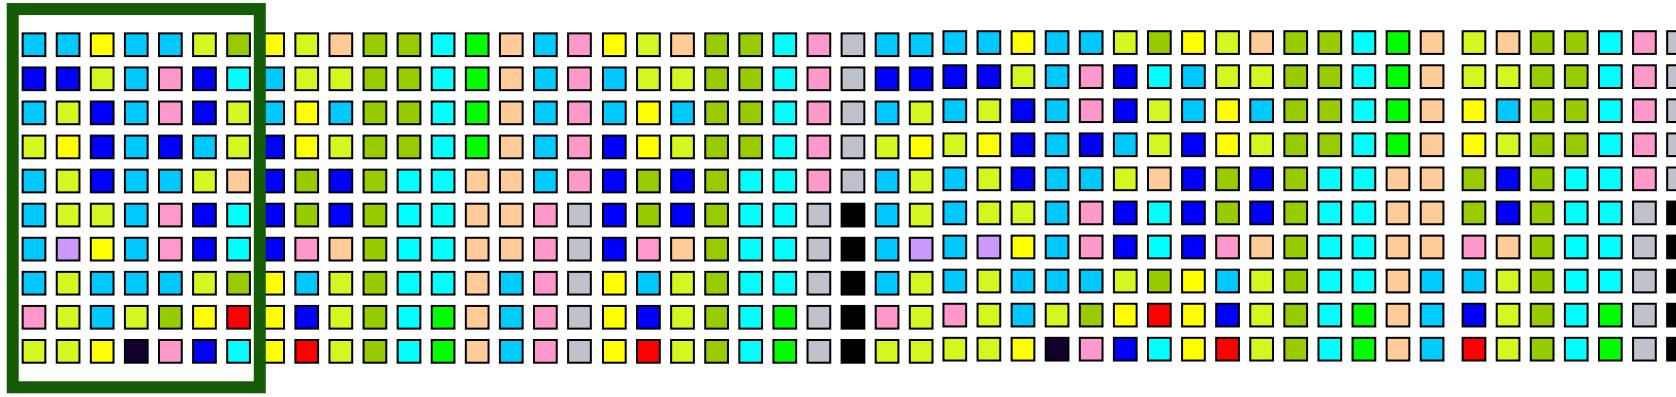
Flux includes a set of window operators that **tumble** across the results of the time range operator

- e.g. `aggregateWindow(every: 60m, fn: mean)`

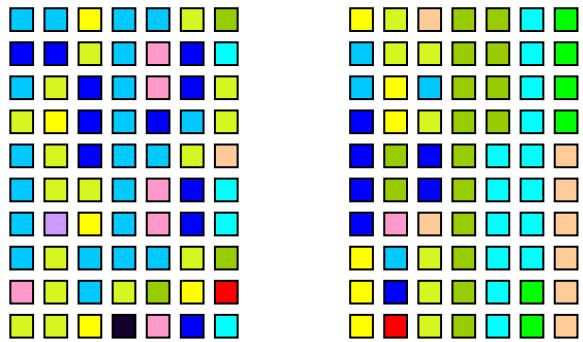
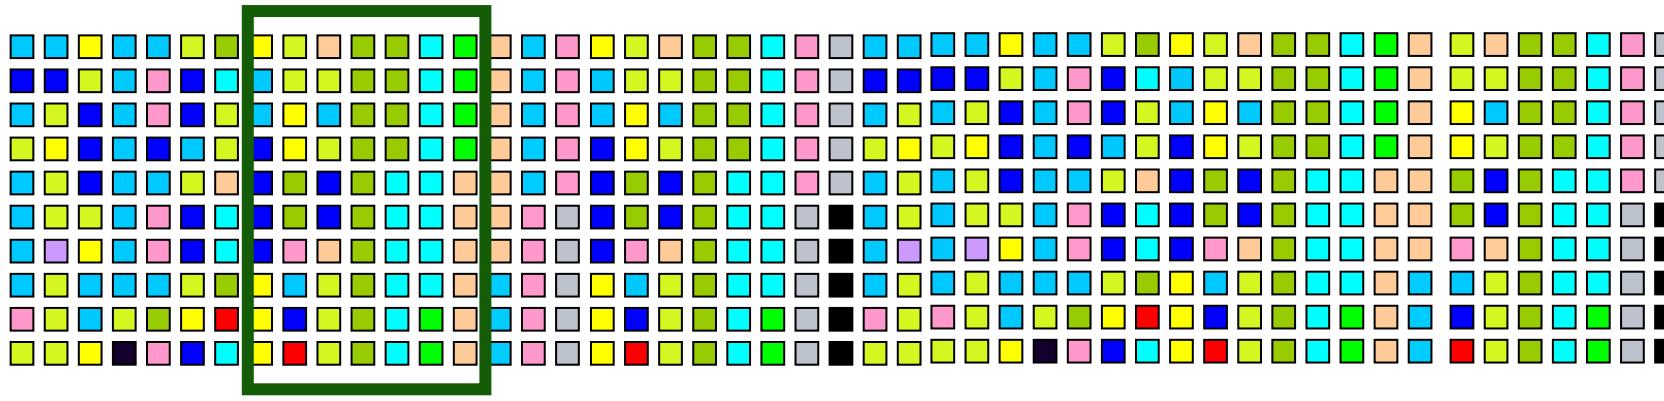
Tumble?



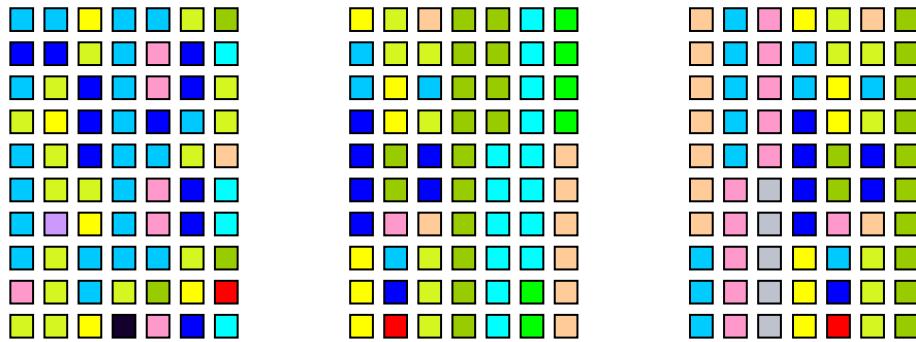
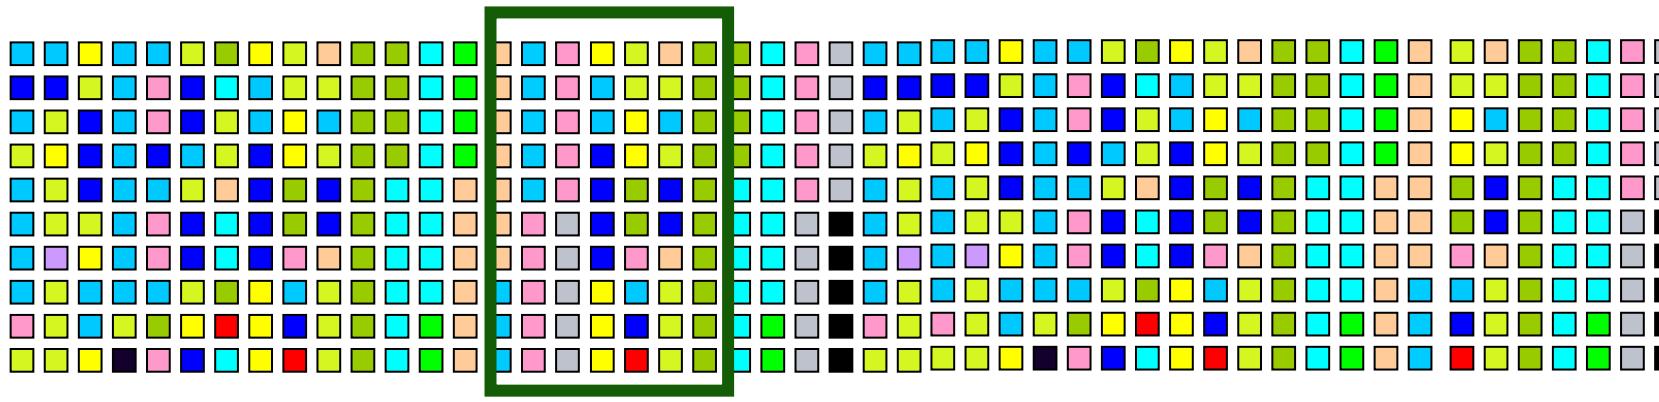
Tumbling windows



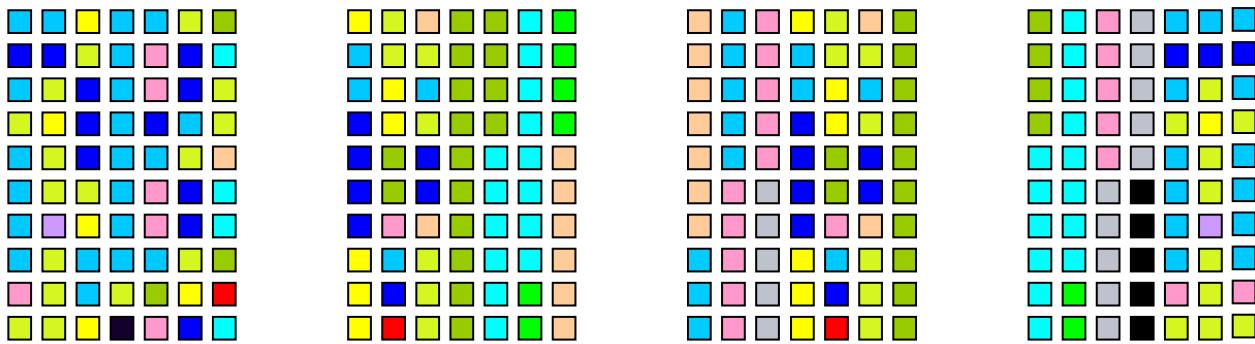
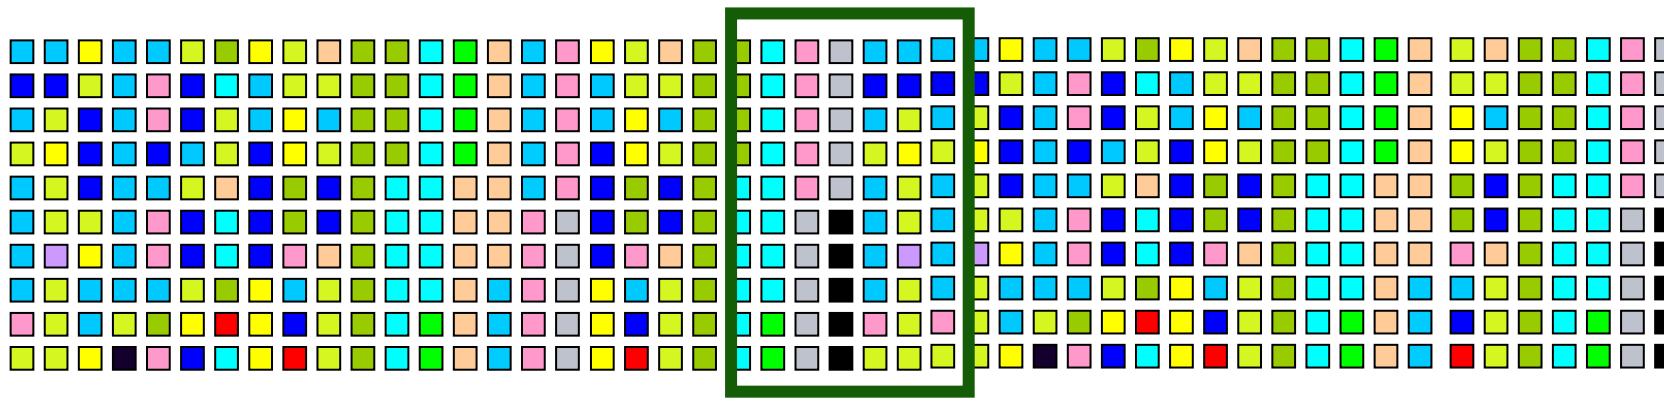
Tumbling windows



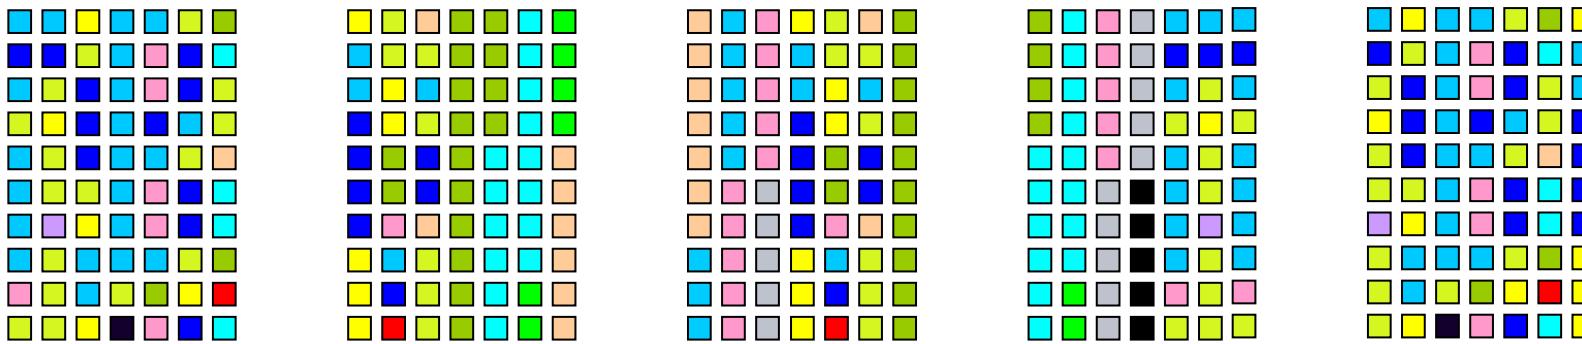
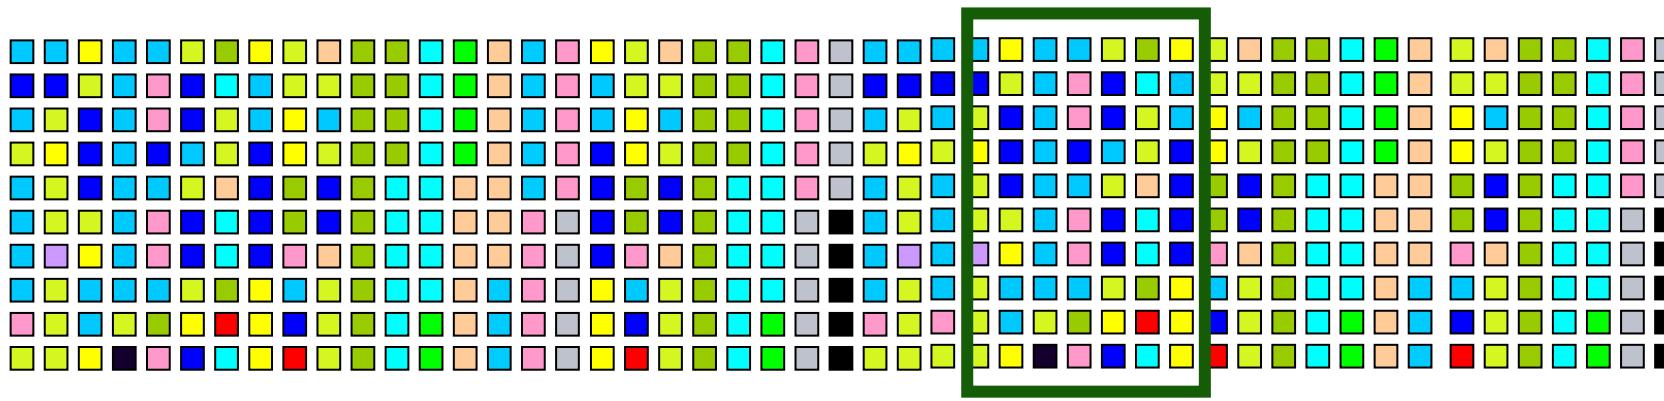
Tumbling windows



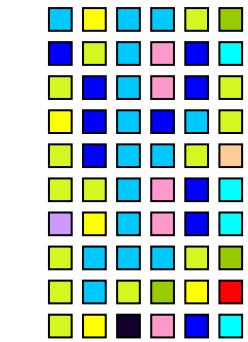
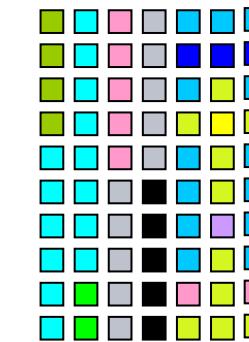
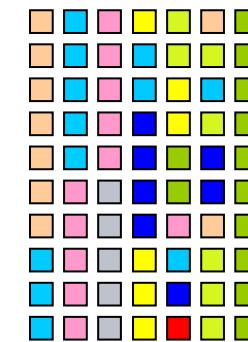
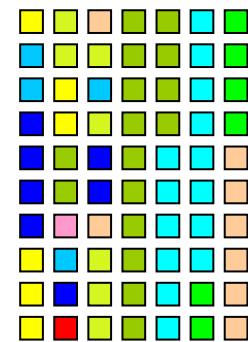
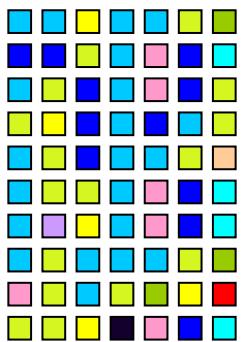
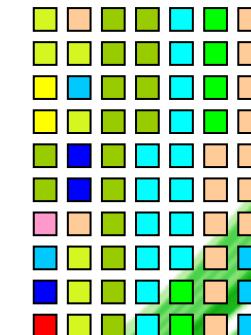
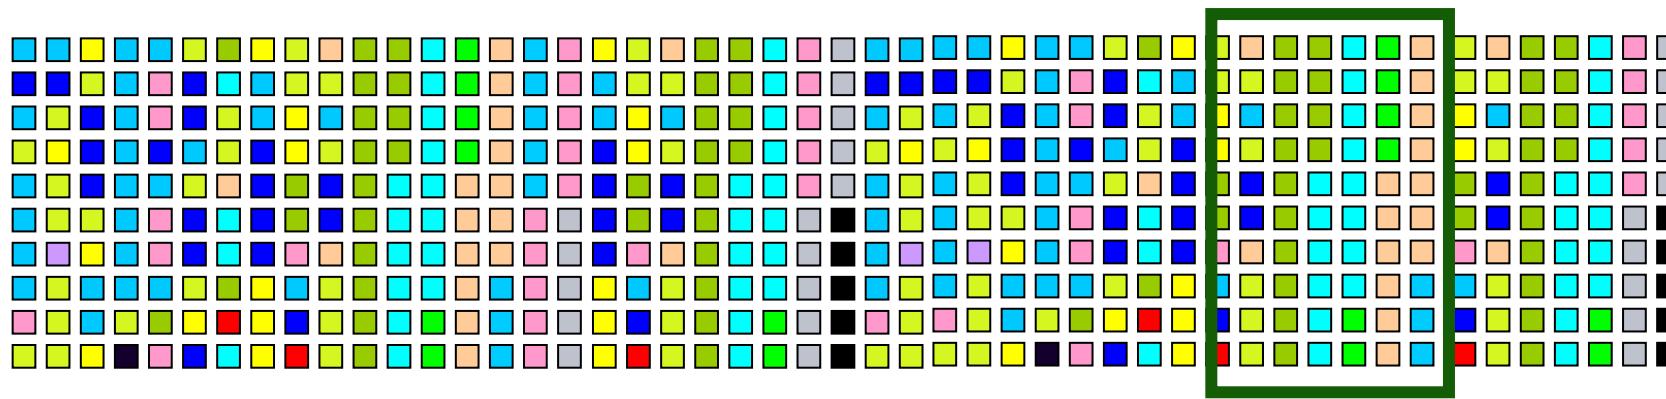
Tumbling windows



Tumbling windows



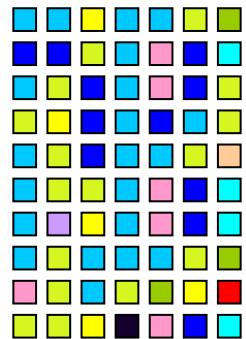
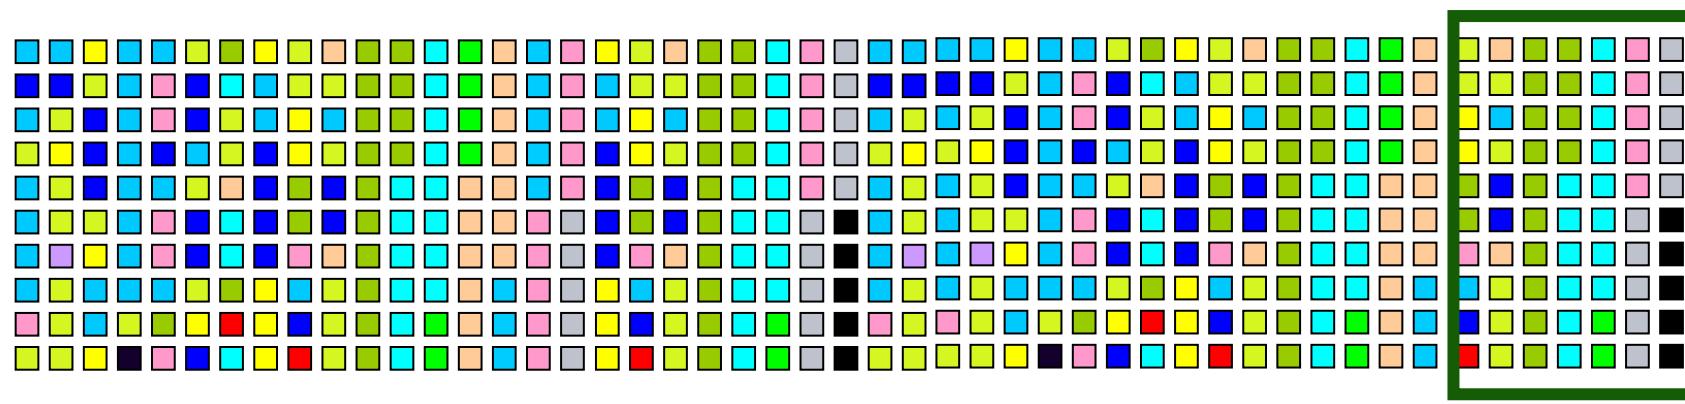
Tumbling windows



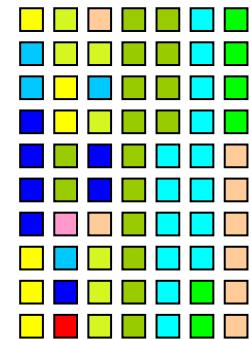
[2]

[6]

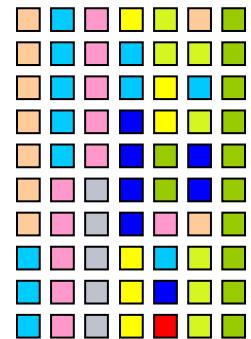
Tumbling windows



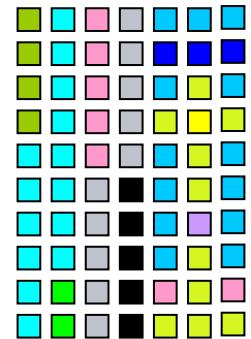
[1]



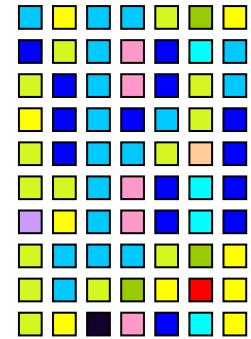
[2]



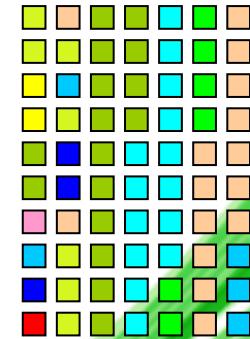
[3]



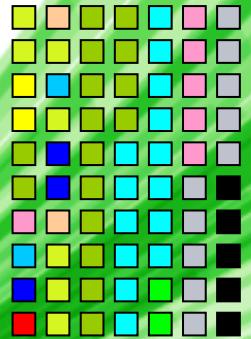
[4]



[5]



[6]



[7]

Recall the Query Process

For Each Table:

1. Convert incoming rows into 0 or more outgoing rows
2. Map outgoing rows to tables

_time	host	_field	_value
10:45	serverA	system	1000
11:00	serverA	system	2000
11:15	serverA	system	1000
11:30	serverA	system	2000
11:45	serverA	system	3000

GroupKey[host=serverA,_field=system]

| >

aggregateWindow(

every: 60m,

in: mean)

_time	host	_field	_value
10:00	serverA	system	1000

GroupKey[host=serverA,_field=system]

_time	host	_field	_value
10:35	serverA	idle	100
11:15	serverA	idle	200
11:50	serverA	idle	400

GroupKey[host=serverA,_field=idle]

_time	host	_field	_value
10:00	serverA	idle	100

GroupKey[host=serverA,_field=idle]

_time	host	_field	_value
10:45	serverA	system	1000
11:00	serverA	system	2000
11:15	serverA	system	1000
11:30	serverA	system	2000
11:45	serverA	system	3000

GroupKey[host=serverA,_field=system]

| >

aggregateWindow(

every: 60m,

in: mean)

_time	host	_field	_value
10:00	serverA	system	1500
11:00	serverA	system	2000

GroupKey[host=serverA,_field=system]

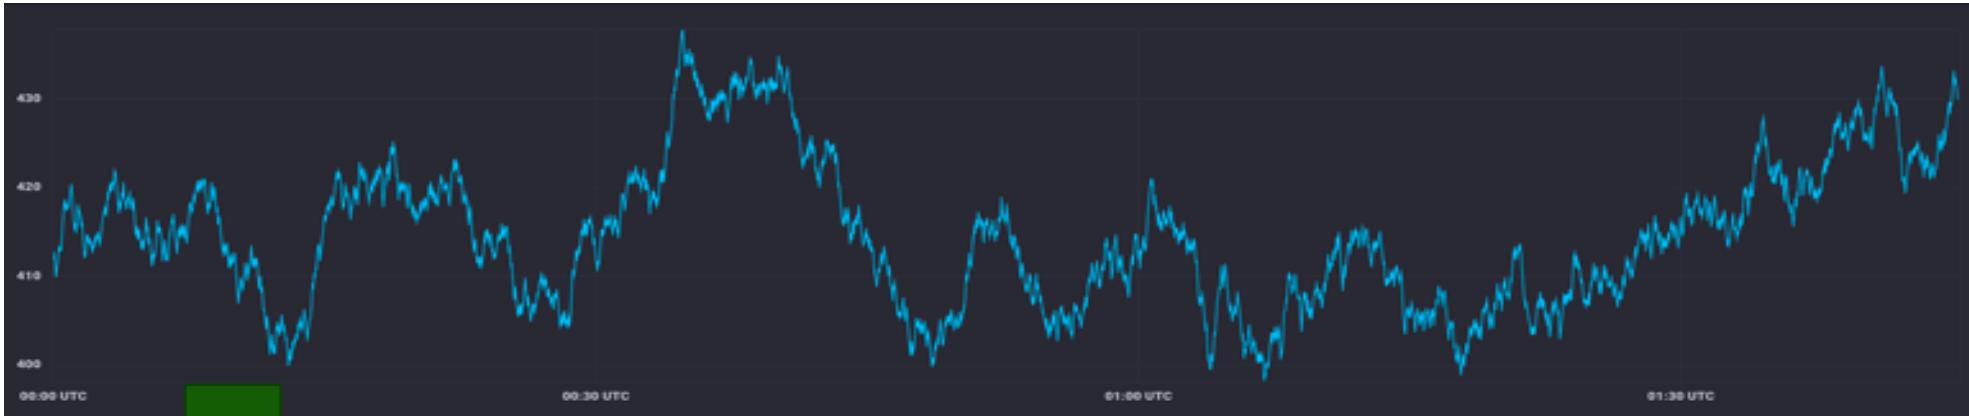
_time	host	_field	_value
10:35	serverA	idle	100
11:15	serverA	idle	200
11:50	serverA	idle	400

GroupKey[host=serverA,_field=idle]

_time	host	_field	_value
10:00	serverA	idle	100
11:00	serverA	idle	300

GroupKey[host=serverA,_field=idle]

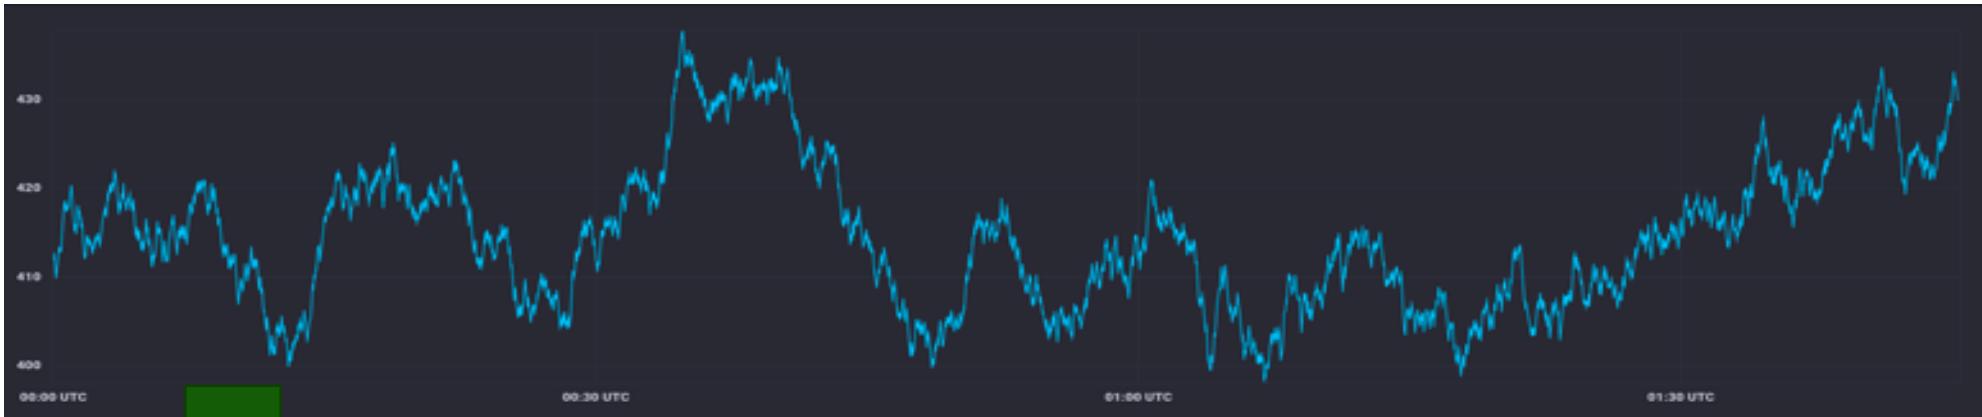
Typical use case: downsampling data



`aggregateWindow(every: 1m, fn: mean)`



Typical use case: downsampling data (cont.)



aggregateWindow(every: 5m, fn: mean)



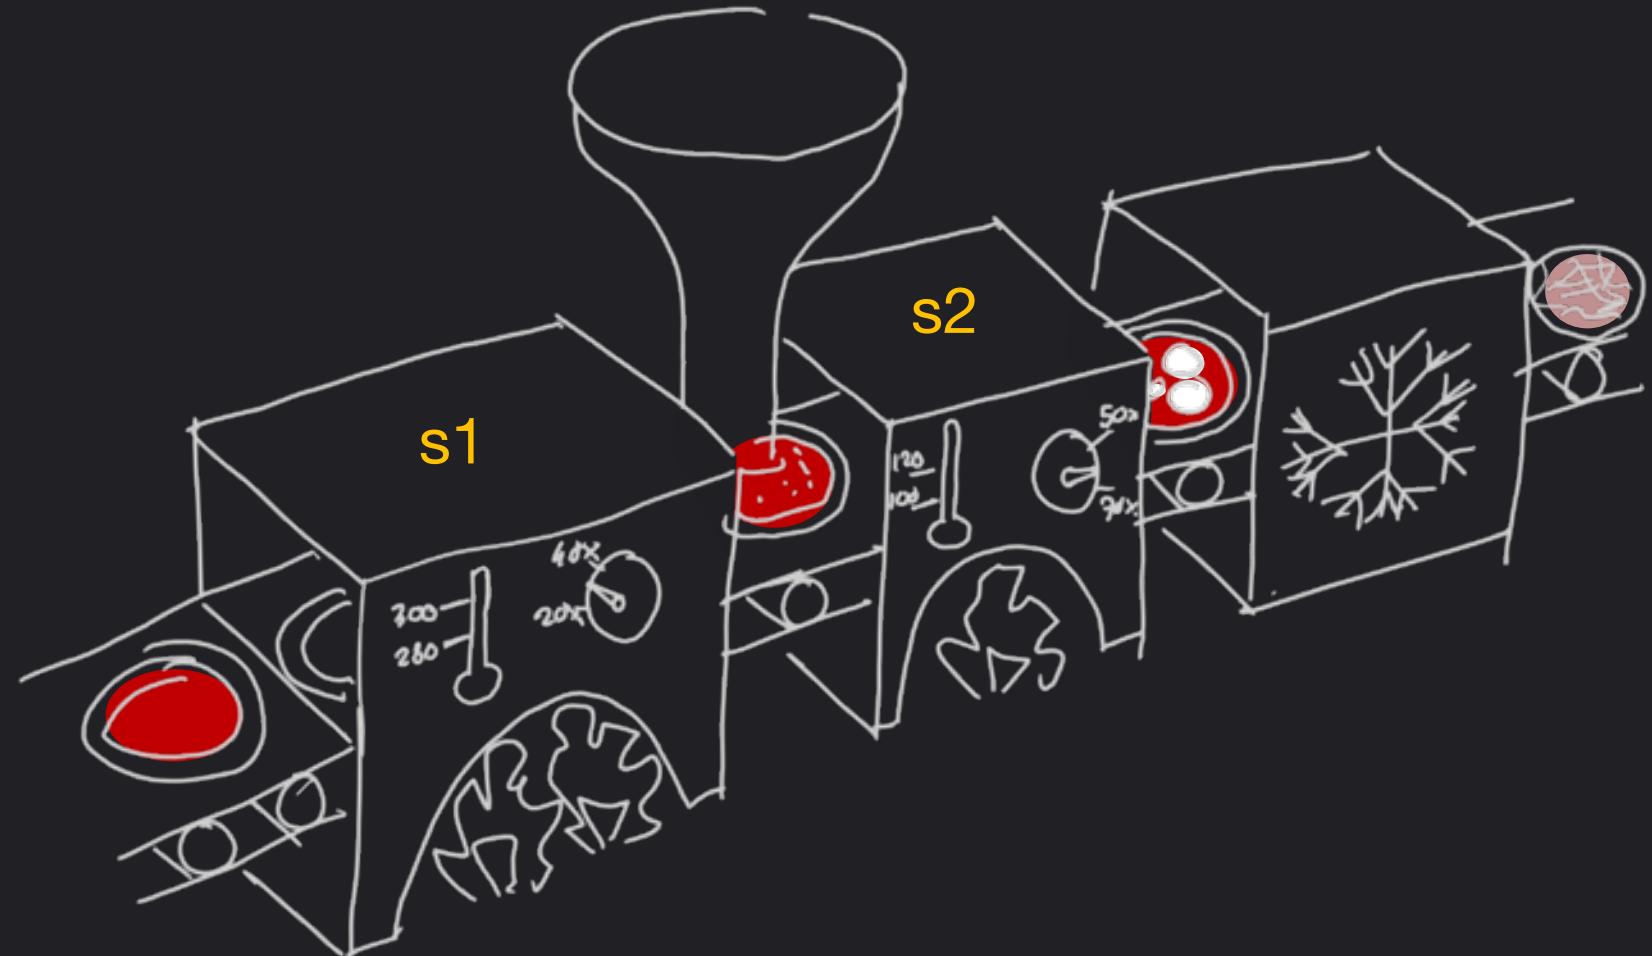
Let's get dirty!



Continuous Linear Pizza Oven

Learning goals:

- aggregateWindow



Task

Extract the moving average of the temperatures observed in the cooking base area over a window of 2 minutes

Extract the moving average of the temperatures observed by S2 over a window of 3 minutes

Take home message

The `aggregateWindow()` function eases the **downsampling** task

- The **createEmpty flag** changes the behaviour of the `aggregateWindow()` function.
 - If `createEmpty = false` the `aggregateWindow()` function does **not** consider the **empty windows** in the aggregation
 - If `createEmpty = true` (default) the `aggregateWindow()` function considers **all the windows**



influx/days